

MIDS W261: Machine Learning at Scale

- **Author:** Karthik Chepudira, Konniam Chan
- **Email:** kchepudira@berkeley.edu, konniam.chan@berkeley.edu
- **Class:** W261-3 Spring 2016
- **Week:** 3
- **Submission Date:** 02-04-2016

HW3.0.

****What is a merge sort? Where is it used in Hadoop? How is a combiner function in the context of Hadoop? Give an example where it can be used and justify why it should be used in the context of this problem. What is the Hadoop shuffle?****

Solution

Merge sort is an algorithm that belongs to family of Merge algorithms that takes an approach of divide and conquer by splitting up a list of n unsorted elements into n sublists with 1 element each and further repeatedly merges sublists while sorting them until only 1 sorted sublist remains.

Conceptually, a merge sort works as follows:

1. Divide the unsorted list into n sublists, each containing 1 element (a list of 1 element is considered sorted).
2. Repeatedly merge sublists to produce new sorted sublists until there is only 1 sublist remaining. This will be the sorted list.

merge sort has a time complexity of $O(n \log n)$

Merge sort in Hadoop could be used in situation where individual sorted lists need to be merged and sorted. This is used in hadoop to merge partition files stream to reducer.

merge sort of sorted lists has a time complexity of $O(n)$

Combiners are an optimization in MapReduce that allows for aggregation before the shuffle phase. In a word count problem in MapReduce combiners and be used for aggregation to reduce the number of tuples transferred across the network from mappers to reducers.

Input:

```
It is a sunny day
It is a rainy day
```

Mapper output :

```
[ (It,1) (is,1) (a,1) (sunny,1) (day,1) (It,1) (is,1) (a,1) (rainy,1)
(day,1) ]
```

Combiner Output:

```
[ (It,2) (is,2) (a,2) (sunny,1) (day,2) (rainy,1) ]
```

Hadoop shuffle refers to processes that happen between the mapper producing the output and the reducer consuming it. When the map function starts producing output it's not simply written to disk. The shuffle partitions, sorts, and combines mapper outputs before sending them to the reducers, taking advantage of buffering write in memory and doing some presorting for efficiency reasons. In the end, the reducer receives a merged partition file sorted by key.


References:

1. [Wikipedia Merge Sort \(https://en.wikipedia.org/wiki/Merge_sort\)](https://en.wikipedia.org/wiki/Merge_sort).
2. [Merge sort in detail \(http://www.personal.kent.edu/~rmuhamma/Algorithms/MyAlgorithms/Sorting/mergeSort.htm\)](http://www.personal.kent.edu/~rmuhamma/Algorithms/MyAlgorithms/Sorting/mergeSort.htm).

```
In [1]: from IPython.display import Image
        Image(filename='HadoopShuffle.png')
```

Out[1]:

MapReduce Overview



HW3.1

Use Counters to do EDA (exploratory data analysis and to monitor progress) Counters are lightweight objects in Hadoop that allow you to keep track of system progress in both the map and reduce stages of processing. By default, Hadoop defines a number of standard counters in "groups"; these show up in the jobtracker webapp, giving you information such as "Map input records", "Map output records", etc.

While processing information/data using MapReduce job, it is a challenge to monitor the progress of parallel threads running across nodes of distributed clusters. Moreover, it is also complicated to distinguish between the data that has been processed and the data which is yet to be processed. The MapReduce Framework offers a provision of user-defined Counters, which can be effectively utilized to monitor the progress of data across nodes of distributed clusters.

Use the Consumer Complaints Dataset provide here to complete this question:

https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints.csv?dl=0

The consumer complaints dataset consists of diverse consumer complaints, which have been reported across the United States regarding various types of loans. The dataset consists of records of the form:

Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?

Here's is the first few lines of the of the Consumer Complaints Dataset:

Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?

1114245,Debt collection,Medical,Disclosure verification of debt,Not given enough info to verify debt,FL,32219,Web,11/13/2014,11/13/2014,"Choice Recovery, Inc.",Closed with explanation,Yes,

1114488,Debt collection,Medical,Disclosure verification of debt,Right to dispute notice not received,TX,75006,Web,11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In progress,Yes,

1114255,Bank account or service,Checking account,Deposits and withdrawals,,NY,11102,Web,11/13/2014,11/13/2014,"FNIS (Fidelity National Information Services, Inc.)",In progress,Yes,

1115106,Debt collection,"Other (phone, health club, etc.)",Communication tactics,Frequent or repeated calls,GA,31721,Web,11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In progress,Yes,

User-defined Counters

Now, let's use Hadoop Counters to identify the number of complaints pertaining to debt collection, mortgage and other categories (all other categories get lumped into this one) in the consumer complaints dataset. Basically produce the distribution of the Product column in this dataset using counters (limited to 3 counters here).

Hadoop offers Job Tracker, an UI tool to determine the status and statistics of all jobs. Using the job tracker UI, developers can view the Counters that have been created. Screenshot your job tracker UI as your job completes and include it here. Make sure that your user defined counters are visible.

```
In [1]: # Start hadoop
!start-dfs.sh
!start-yarn.sh
```

```
16/01/29 11:22:41 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /Users/InfernoIX/hadoop-2.
7.1/logs/hadoop-InfernoIX-namenode-Konniam-MacBook-Air.local.out
localhost: starting datanode, logging to /Users/InfernoIX/hadoop-2.
7.1/logs/hadoop-InfernoIX-datanode-Konniam-MacBook-Air.local.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /Users/InfernoIX/ha
doo-2.7.1/logs/hadoop-InfernoIX-secondarynamenode-Konniam-MacBook
-Air.local.out
16/01/29 11:23:03 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
starting yarn daemons
starting resourcemanager, logging to /Users/InfernoIX/hadoop-2.7.1/
logs/yarn-InfernoIX-resourcemanager-Konniam-MacBook-Air.local.out
localhost: starting nodemanager, logging to /Users/InfernoIX/hadoop
-2.7.1/logs/yarn-InfernoIX-nodemanager-Konniam-MacBook-Air.local.o
ut
```

```
In [2]: # Remove header of the CSV file to simplify later steps
# !tail -n +2 Consumer_Complaints.csv > Consumer_Fixed.csv
# !mv Consumer_Fixed.csv Consumer_Complaints.csv
```

```
In [7]: # Place in HDFS
!hdfs dfs -put Consumer_Complaints.csv /user/konniam/week_03
```

```
16/01/29 12:04:06 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

```
In [3]: %%writefile mapper.py
#!/usr/bin/env python
# Use mapper to do EDA using counters
import sys

for line in sys.stdin:
    complaint = line.strip().split(',')
    # Second column contains the product info
    if complaint[1] == "Debt collection":
        sys.stderr.write("reporter:counter:Debt-Counter,debt,1\n")
    elif complaint[1] == "Mortgage":
        sys.stderr.write("reporter:counter:Mortgage-Counter,mortgage,1\n")
    else:
        sys.stderr.write("reporter:counter:Other-Counter,other,1\n")
    # Output original line
    print line
```

Writing mapper.py

```
In [4]: !chmod a+x mapper.py
```

```
In [39]: !hdfs dfs -rm -r /user/konniam/week_03/hw_3_1_output
```

```
16/01/29 15:03:36 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
16/01/29 15:03:39 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /user/konniam/week_03/hw_3_1_output
```

```
In [5]: !hdfs dfs -rm -r /user/konniam/week_03/hw_3_1_output
# Run MR job to increment counters, with Identity Reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.jar \
-mapper $PWD/mapper.py \
-reducer /bin/cat \
-input /user/konniam/week_03/Consumer_Complaints.csv \
-output /user/konniam/week_03/hw_3_1_output
```

```
16/01/31 16:37:59 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
16/01/31 16:38:00 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /user/konniam/week_03/hw_3_1_output
```

```

16/01/31 16:38:02 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h5l_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar5873106299045917131/] [] /var/folders/l8/h5l_59852qscq403
fs6q0xlh0000gn/T/streamjob1911821617765572003.jar tmpDir=null
16/01/31 16:38:03 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/01/31 16:38:04 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/01/31 16:38:05 INFO mapred.FileInputFormat: Total input paths to
process : 1
16/01/31 16:38:05 INFO mapreduce.JobSubmitter: number of splits:2
16/01/31 16:38:05 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454277284610_0002
16/01/31 16:38:06 INFO impl.YarnClientImpl: Submitted application a
pplication_1454277284610_0002
16/01/31 16:38:06 INFO mapreduce.Job: The url to track the job: htt
p://Konniams-MacBook-Air.local:8088/proxy/application_1454277284610
_0002/
16/01/31 16:38:06 INFO mapreduce.Job: Running job: job_145427728461
0_0002
16/01/31 16:38:16 INFO mapreduce.Job: Job job_1454277284610_0002 ru
nning in uber mode : false
16/01/31 16:38:16 INFO mapreduce.Job:  map 0% reduce 0%
16/01/31 16:38:32 INFO mapreduce.Job:  map 67% reduce 0%
16/01/31 16:38:33 INFO mapreduce.Job:  map 100% reduce 0%
16/01/31 16:38:45 INFO mapreduce.Job:  map 100% reduce 100%
16/01/31 16:38:45 INFO mapreduce.Job: Job job_1454277284610_0002 co
mpleted successfully
16/01/31 16:38:45 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=53659840
        FILE: Number of bytes written=107672125
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50910032
        HDFS: Number of bytes written=51845048
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=30890
        Total time spent by all reduces in occupied slots (
ms)=8230

```

```

Total time spent by all map tasks (ms)=30890
Total time spent by all reduce tasks (ms)=8230
Total vcore-seconds taken by all map tasks=30890
Total vcore-seconds taken by all reduce tasks=8230
Total megabyte-seconds taken by all map tasks=31631
360
Total megabyte-seconds taken by all reduce tasks=84
27520

```

Map-Reduce Framework

```

Map input records=312912
Map output records=625824
Map output bytes=52126023
Map output materialized bytes=53659846
Input split bytes=236
Combine input records=0
Combine output records=0
Reduce input groups=312913
Reduce shuffle bytes=53659846
Reduce input records=625824
Reduce output records=625824
Spilled Records=1251648
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=819
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=587726848

```

Debt-Counter

```
debt=44372
```

Mortgage-Counter

```
mortgage=125752
```

Other-Counter

```
other=142788
```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=50909796
```

File Output Format Counters

```
Bytes Written=51845048
```

```

16/01/31 16:38:45 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_1_output

```



```
In [7]: from IPython.display import Image
```

3.1 Response

Counters shown below:

```
In [9]: Image(filename='hw_3_1_counter.png')
```

Out[9]:

	Name	Map	Reduce	Total
Map-Reduce Framework	Combine input records	0	0	0
	Combine output records	0	0	0
	CPU time spent (ms)	0	0	0
	Failed Shuffles	0	0	0
	GC time elapsed (ms)	510	309	819
	Input split bytes	236	0	236
	Map input records	312,912	0	312,912
	Map output bytes	52,126,023	0	52,126,023
	Map output materialized bytes	53,659,846	0	53,659,846
	Map output records	625,824	0	625,824
	Merged Map outputs	0	2	2
	Physical memory (bytes) snapshot	0	0	0
	Reduce input groups	0	312,913	312,913
	Reduce input records	0	625,824	625,824
	Reduce output records	0	625,824	625,824
	Reduce shuffle bytes	0	53,659,846	53,659,846
	Shuffled Maps	0	2	2
	Spilled Records	625,824	625,824	1,251,648
	Total committed heap usage (bytes)	393,740,288	193,986,560	587,726,848
	Virtual memory (bytes) snapshot	0	0	0
Debt-Counter	debt	44,372	0	44,372
Mortgage-Counter	mortgage	125,752	0	125,752
Other-Counter	other	142,788	0	142,788
Shuffle Errors	BAD_ID	0	0	0
	CONNECTION	0	0	0
	IO_ERROR	0	0	0
	WRONG_LENGTH	0	0	0
	WRONG_MAP	0	0	0
File Input Format Counters	Bytes Read	50,909,796	0	50,909,796
	Bytes Written	0	51,845,048	51,845,048

HW 3.2 Analyze the performance of your Mappers, Combiners and Reducers using Counters

3.2 Part 1

For this brief study the Input file will be one record (the next line only): foo foo quux labs foo bar quux

Perform a word count analysis of this single record dataset using a Mapper and Reducer based WordCount (i.e., no combiners are used here) using user defined Counters to count up how many time the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing this word count job. The answer should be 1 and 4 respectively. Please explain.

```
In [4]: !echo "foo foo quux labs foo bar quux" > input3.2.txt
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_2.txt
!hdfs dfs -put input3.2.txt /user/konniam/week_03/hw_3_2.txt
!hdfs dfs -cat /user/konniam/week_03/hw_3_2.txt
```

16/02/01 03:13:17 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.

Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/input3.2.txt' to trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Current

foo foo quux labs foo bar quux

Mapper

```
In [1]: %%writefile mapper3.2.py
#!/usr/bin/env python
## reducer3.2.py
## Author: karthik chepudira
## Description: mapper code for HW3.2
import sys

sys.stderr.write("reporter:counter:MRCounter,Mapper,1\n") #increment the counter for Mapper
for line in sys.stdin:
    if not line in ['\n', '\r\n']: # check for empty lines and ignore
        words = line.rstrip().split()
        for word in words:
            print word+'\t'+str(1)
```

Overwriting mapper3.2.py

Reducer

```
In [2]: %%writefile reducer3.2.py
#!/usr/bin/env python
## reducer3.2.py
## Author: karthik chepudira
## Description: reducer code for HW3.2

import sys
sys.stderr.write("reporter:counter:MRCounter,Reducer,1\n") # increment the counter for Reducer

#function to print reducer output by key
def wcount(prev_word ,counts):
    if prev_word != "-1": # check if a valid key
        print(prev_word+ "\t"+str(counts)) # Print reducer output

prev_word = "-1"
counts = 0
linecount=0 #Keep track of total lines of input processed by reducer
for line in sys.stdin:
    line = line.rstrip()
    linecount+=1
    word, value =line.split("\t",1)
    if word!=prev_word: #if all rows of a key are processed
        wcount(prev_word, counts) # call the print function
        prev_word = word # reset key
        counts = 0 #reset key count
    counts += eval(value)
if linecount >0: # if reducer processed atleast 1 line
    if word==prev_word:
        wcount(prev_word, counts)
```

Overwriting reducer3.2.py

```
In [3]: !chmod a+x mapper3.2.py reducer3.2.py
```

```
In [6]: !hdfs dfs -rm -r /tmp/3.2/streamout
# Run word count to increment mapper and reducer counters
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2" \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.map.tasks=1 \
-D mapred.reduce.tasks=4 \
-input /tmp/3.2/input3.2.txt \
-output /tmp/3.2/streamout \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2.py' \
```

16/02/02 04:07:25 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.

Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/streamout' to trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Current

WARNING: Use "yarn jar" to launch YARN applications.

packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob9161353458867003514.jar tmpDir=null

16/02/02 04:07:33 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/02 04:07:33 INFO client.RMProxy: Connecting to ResourceManager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/02 04:07:35 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/02 04:07:35 INFO client.RMProxy: Connecting to ResourceManager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/02 04:07:36 INFO mapred.FileInputFormat: Total input paths to process : 1

16/02/02 04:07:37 INFO mapreduce.JobSubmitter: number of splits:1

16/02/02 04:07:39 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454374677022_0001

16/02/02 04:07:41 INFO impl.YarnClientImpl: Submitted application application_1454374677022_0001

16/02/02 04:07:42 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454374677022_0001/

16/02/02 04:07:42 INFO mapreduce.Job: Running job: job_1454374677022_0001

16/02/02 04:08:05 INFO mapreduce.Job: Job job_1454374677022_0001 running in uber mode : false

16/02/02 04:08:05 INFO mapreduce.Job: map 0% reduce 0%

16/02/02 04:08:13 INFO mapreduce.Job: map 100% reduce 0%

16/02/02 04:08:25 INFO mapreduce.Job: map 100% reduce 25%

```

16/02/02 04:08:27 INFO mapreduce.Job:  map 100% reduce 50%
16/02/02 04:08:28 INFO mapreduce.Job:  map 100% reduce 75%
16/02/02 04:08:29 INFO mapreduce.Job:  map 100% reduce 100%
16/02/02 04:08:31 INFO mapreduce.Job:  Job job_1454374677022_0001 completed successfully

```

```

16/02/02 04:08:32 INFO mapreduce.Job: Counters: 51

```

File System Counters

```

FILE: Number of bytes read=83
FILE: Number of bytes written=639228
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=139
HDFS: Number of bytes written=26
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=8

```

Job Counters

```

Launched map tasks=1
Launched reduce tasks=4
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)
=6795
Total time spent by all reduces in occupied slots (
ms)=41630
Total time spent by all map tasks (ms)=6795
Total time spent by all reduce tasks (ms)=41630
Total vcore-seconds taken by all map tasks=6795
Total vcore-seconds taken by all reduce tasks=41630
Total megabyte-seconds taken by all map tasks=16987
50
Total megabyte-seconds taken by all reduce tasks=10
407500

```

Map-Reduce Framework

```

Map input records=1
Map output records=7
Map output bytes=45
Map output materialized bytes=83
Input split bytes=108
Combine input records=0
Combine output records=0
Reduce input groups=4
Reduce shuffle bytes=83
Reduce input records=7
Reduce output records=4
Spilled Records=14
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=891

```

```
CPU time spent (ms)=7460
Physical memory (bytes) snapshot=762757120
Virtual memory (bytes) snapshot=5174779904
Total committed heap usage (bytes)=622854144
MRCounter
  Mapper=1
  Reducer=4
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=31
File Output Format Counters
  Bytes Written=26
16/02/02 04:08:32 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout
```

3.2 Part 1 Response

1. Invoke the hadoop job specifying the number of Map tasks and reduce tasks by using the following options
 - -D mapred.map.tasks=1 \
 - -D mapred.reduce.tasks=4 \

3.2 Part 2

Please use multiple mappers and reducers for these jobs (at least 2 mappers and 2 reducers).

Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper and Reducer based WordCount (i.e., no combiners used anywhere) using user defined Counters to count up how many time the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

Mapper 3.2 Part 2

```
In [7]: %%writefile mapper3.2_2.py
#!/usr/bin/env python
## reducer3.2_2.py
## Author: karthik chepudira
## Description: mapper code for HW3.2 part 2
import sys
import re
from nltk.corpus import stopwords #stopwords from nltk
from csv import reader #csv reader

sys.stderr.write("reporter:counter:MRCCounter,Mapper,1\n") #increment
the counter for Mapper

# function to remove punctuation from string
def removepunctuation(inputstring):
    punctuations = '''!()-[]{};:'"\,<>./?@#$$%^&*~' '''
    ret_string = ""
    for char in inputstring:
        if char not in punctuations:
            ret_string = ret_string + char
    return ret_string

total_count = 0
linecount=0 #Keep track of total lines of input processed by mapper
for line in reader(sys.stdin): # use csvreader to process input lin
e.
    if not line in ['\n', '\r\n']:
        linecount+=1
        columns = line
        colissue=columns[3] #third column of csv file has the 'issu
e' field
        #remove punctuation , lower , split and check alphanumeric.
        words = [w for w in removepunctuation(colissue).lower().spl
it() if not w in stopwords.words('english') and w.isalpha()]
        for word in words:
            print word+'\t'+str(1)
            total_count+=1
if linecount > 0: # if mapper processed atleast 1 line
    print "*"+"\t"+str(total_count)
```

Overwriting mapper3.2_2.py

Reducer 3.2 Part 2

```
In [8]: %%writefile reducer3.2_2.py
#!/usr/bin/env python
## reducer3.2_2.py
## Author: karthik chepudira
## Description: reducer code for HW3.2 part 2
import sys
prev_key="-1"
count=0
total_count=0 #keep count of the keys processed
for line in sys.stdin:
    key,value=line.split("\t",1)
    if key == "*":
        total_count+= eval(value) #increment the total key count from the mapper output.
    else:
        if prev_key != key:
            if prev_key != "-1":
                print prev_key+"\t"+str(count) #print key,value pair
                count = 0 #reset count
            prev_key = key # reset key
            count+=eval(value) #increment count
        else:
            count+=eval(value) # increment count
print prev_key+"\t"+str(count) # print the last key
print "*"+"\t"+str(total_count) #print the total key count.
sys.stderr.write("reporter:counter:MRCOUNTER,Reducer,1\n") # increment the counter for Reducer
```

Overwriting reducer3.2_2.py

```
In [9]: !chmod a+x mapper3.2_2.py reducer3.2_2.py
```

```
In [2]: #remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout2
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2_2" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=2 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.text.key.comparator.options="-k1" \
-input /tmp/3.1/Consumer_Complaints.csv \
-output /tmp/3.2/streamout2 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_2.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_2.py'

rm: `/tmp/3.2/streamout2': No such file or directory
```



```
WARNING: Use "yarn jar" to launch YARN applications.
packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob8758615865896012356.jar tmpDir=null
16/02/02 20:58:21 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/02 20:58:21 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050
16/02/02 20:58:22 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/02 20:58:22 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050
16/02/02 20:58:23 INFO mapred.FileInputFormat: Total input paths to process : 1
16/02/02 20:58:23 INFO mapreduce.JobSubmitter: number of splits:2
16/02/02 20:58:24 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454426531297_0003
16/02/02 20:58:24 INFO impl.YarnClientImpl: Submitted application application_1454426531297_0003
16/02/02 20:58:25 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454426531297_0003/
16/02/02 20:58:25 INFO mapreduce.Job: Running job: job_1454426531297_0003
16/02/02 20:58:32 INFO mapreduce.Job: Job job_1454426531297_0003 running in uber mode : false
16/02/02 20:58:32 INFO mapreduce.Job: map 0% reduce 0%
16/02/02 20:58:44 INFO mapreduce.Job: map 3% reduce 0%
16/02/02 20:58:47 INFO mapreduce.Job: map 5% reduce 0%
16/02/02 20:58:50 INFO mapreduce.Job: map 7% reduce 0%
16/02/02 20:58:53 INFO mapreduce.Job: map 9% reduce 0%
16/02/02 20:58:56 INFO mapreduce.Job: map 11% reduce 0%
16/02/02 20:58:59 INFO mapreduce.Job: map 14% reduce 0%
16/02/02 20:59:02 INFO mapreduce.Job: map 15% reduce 0%
16/02/02 20:59:03 INFO mapreduce.Job: map 16% reduce 0%
16/02/02 20:59:05 INFO mapreduce.Job: map 17% reduce 0%
16/02/02 20:59:06 INFO mapreduce.Job: map 18% reduce 0%
16/02/02 20:59:08 INFO mapreduce.Job: map 19% reduce 0%
16/02/02 20:59:09 INFO mapreduce.Job: map 20% reduce 0%
16/02/02 20:59:11 INFO mapreduce.Job: map 21% reduce 0%
16/02/02 20:59:12 INFO mapreduce.Job: map 22% reduce 0%
16/02/02 20:59:14 INFO mapreduce.Job: map 23% reduce 0%
16/02/02 20:59:15 INFO mapreduce.Job: map 24% reduce 0%
16/02/02 20:59:17 INFO mapreduce.Job: map 25% reduce 0%
16/02/02 20:59:18 INFO mapreduce.Job: map 26% reduce 0%
16/02/02 20:59:20 INFO mapreduce.Job: map 27% reduce 0%
16/02/02 20:59:21 INFO mapreduce.Job: map 28% reduce 0%
16/02/02 20:59:24 INFO mapreduce.Job: map 29% reduce 0%
16/02/02 20:59:27 INFO mapreduce.Job: map 30% reduce 0%
16/02/02 20:59:30 INFO mapreduce.Job: map 32% reduce 0%
```

```

16/02/02 20:59:33 INFO mapreduce.Job: map 33% reduce 0%
16/02/02 20:59:36 INFO mapreduce.Job: map 35% reduce 0%
16/02/02 20:59:39 INFO mapreduce.Job: map 37% reduce 0%
16/02/02 20:59:42 INFO mapreduce.Job: map 39% reduce 0%
16/02/02 20:59:45 INFO mapreduce.Job: map 41% reduce 0%
16/02/02 20:59:48 INFO mapreduce.Job: map 43% reduce 0%
16/02/02 20:59:51 INFO mapreduce.Job: map 45% reduce 0%
16/02/02 20:59:54 INFO mapreduce.Job: map 47% reduce 0%
16/02/02 20:59:57 INFO mapreduce.Job: map 49% reduce 0%
16/02/02 21:00:00 INFO mapreduce.Job: map 51% reduce 0%
16/02/02 21:00:03 INFO mapreduce.Job: map 53% reduce 0%
16/02/02 21:00:06 INFO mapreduce.Job: map 56% reduce 0%
16/02/02 21:00:09 INFO mapreduce.Job: map 58% reduce 0%
16/02/02 21:00:12 INFO mapreduce.Job: map 60% reduce 0%
16/02/02 21:00:15 INFO mapreduce.Job: map 63% reduce 0%
16/02/02 21:00:18 INFO mapreduce.Job: map 81% reduce 0%
16/02/02 21:00:24 INFO mapreduce.Job: map 82% reduce 0%
16/02/02 21:00:30 INFO mapreduce.Job: map 83% reduce 0%
16/02/02 21:00:34 INFO mapreduce.Job: map 83% reduce 8%
16/02/02 21:00:35 INFO mapreduce.Job: map 100% reduce 17%
16/02/02 21:00:38 INFO mapreduce.Job: map 100% reduce 77%
16/02/02 21:00:41 INFO mapreduce.Job: map 100% reduce 94%
16/02/02 21:00:42 INFO mapreduce.Job: map 100% reduce 100%
16/02/02 21:00:42 INFO mapreduce.Job: Job job_1454426531297_0003 completed successfully
16/02/02 21:00:42 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=13877162
        FILE: Number of bytes written=28266186
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=51012846
        HDFS: Number of bytes written=2032
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=224050
        Total time spent by all reduces in occupied slots (
ms)=41406
        Total time spent by all map tasks (ms)=224050
        Total time spent by all reduce tasks (ms)=41406
        Total vcore-seconds taken by all map tasks=224050
        Total vcore-seconds taken by all reduce tasks=41406
        Total megabyte-seconds taken by all map tasks=56012

```

500

Total megabyte-seconds taken by all reduce tasks=10

351500

Map-Reduce Framework

```

Map input records=312919
Map output records=1002221
Map output bytes=11872708
Map output materialized bytes=13877174
Input split bytes=238
Combine input records=0
Combine output records=0
Reduce input groups=152
Reduce shuffle bytes=13877174
Reduce input records=1002221
Reduce output records=153
Spilled Records=2004442
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=439
CPU time spent (ms)=216790
Physical memory (bytes) snapshot=693493760
Virtual memory (bytes) snapshot=4098723840
Total committed heap usage (bytes)=506986496

```

MRCounter

```

Mapper=2
Reducer=2

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

Bytes Read=51012608

File Output Format Counters

Bytes Written=2032

```

16/02/02 21:00:42 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout2

```

```

In [8]: !hdfs dfs -cat /tmp/3.2/streamout2/part* | sort -k2,2rn > output_3.
2_2.txt
!cat output_3.2_2.txt

```

```

*          1002219
loan       119491
modificationcollectionforeclosure      70487
account 57448
credit  50897

```

payments	39993
escrow 36767	
servicing	36767
report 30546	
incorrect	29069
information	29069
debt 27876	
attempts	17972
collect 17972	
contd 17972	
owed 17972	
closing 16205	
management	16205
opening 16205	
deposits	10555
withdrawals	10555
problems	9484
application	8868
communication	8671
tactics 8671	
broker 8625	
mortgage	8625
originator	8625
unable 8178	
billing 8158	
disclosure	7655
verification	7655
disputes	6938
reporting	6560
lease 6337	
caused 5663	
funds 5663	
low 5663	
process 5505	
managing	5006
improper	4966
companys	4858
investigation	4858
identity	4729
card 4407	
get 4357	
reportcredit	4357
score 4357	
costs 4350	
settlement	4350
interest	4238
protection	4143
repaying	3844
fraud 3842	
pay 3821	

contact	3710	
money	3639	
false	3621	
representation	3621	
statements	3621	
info	3553	
sharing	3489	
apr	3431	
rate	3431	
embezzlement	3276	
theft	3276	
makingreceiving	3226	
sending	3226	
fee	3198	
action	2964	
illegal	2964	
takingthreatening		2964
closingcancelling		2795
decision	2774	
underwriting	2774	
customer	2734	
atm	2422	
debit	2422	
using	2422	
lender	2165	
cant	1999	
dealing	1944	
servicer		1944
collection		1907
late	1797	
line	1732	
repay	1647	
service	1518	
determination		1490
transaction		1486
use	1477	
monitoring		1453
relations		1367
taking	1242	
statement		1220
advertising		1193
marketing		1193
payoff	1155	
increasedecrease		1149
issue	1100	
delinquent		1061
practices		1003
rewards	1002	
charged	976	
didnt	925	

dispute	904	
expect	807	
fees	807	
shopping		672
issuance		640
unsolicited		640
balance	598	
transfer		598
scam	566	
issues	538	
changes	350	
forbearance		350
plans	350	
terms	350	
workout	350	
getting	291	
available		274
promised		274
delay	243	
processing		243
advance	240	
cash	240	
privacy	240	
bankruptcy		222
received		216
bank	202	
wrong	169	
arbitration		168
acct	163	
applied	139	
loandid	139	
receive	139	
sale	139	
charges	131	
stop	131	
overlimit		127
apply	118	
amount	98	
credited		92
payment	92	
checks	75	
convenience		75
amt	71	
day	71	
disclosures		64
incorrectmissing		64
citibank	1	
*	0	

```
In [1]: %%writefile wordcountanalysis.py
#!/usr/bin/env python
## wordcountanalysis.py
## Author: karthik chepudira
## Description: HW3.2 part 2
import matplotlib.pyplot as plt
from matplotlib.font_manager import FontProperties
import pandas as pd
import numpy as np
import sys

plt.switch_backend('agg')


def plotmodel(filename,plotname):
    df = pd.read_csv(filename,sep='\t',header=0)
    df.columns =['Word','Count']
    df['Count'] = df['Count'].astype('int')
    fig, ax = plt.subplots()
    fig.suptitle('Histogram for Wordcount')
    df.head(10).plot(x='Word', y='Count',ax = ax, kind='bar')
    ax.set_ylabel('Frequency')
    ax.set_xlabel('Top 10 Words')
    fontP = FontProperties()
    fontP.set_size('small')
    fig.savefig(plotname+".jpg",bbox_inches='tight')

if len(sys.argv) > 2:
    plotname = str(sys.argv[2])
    filename = str(sys.argv[1])
    plotmodel(filename,plotname)
else:
    print "Incorrect number of parameters: please provide model filename "
    exit()
```

Overwriting wordcountanalysis.py

```
In [2]: !chmod a+x wordcountanalysis.py
! ./wordcountanalysis.py "output_3.2_2.txt" 'HW3.2'
from IPython.display import Image
Image(filename='HW3.2.jpg')
```

Out[2]:



3.2 Part 2 Response

The value of the user defined Mapper Counter is: 2

The value of the user defined Reducer Counter is: 2

```
MRCounter
Mapper=2
Reducer=2
```

3.2 Part 3

Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper, Reducer, and standalone combiner (i.e., not an in-memory combiner) based WordCount using user defined Counters to count up how many time the mapper, combiner, reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

Combiner 3.2_3.py

```
In [10]: %%writefile combiner3.2_3.py
#!/usr/bin/env python
## combiner3.2_3.py
## Author: karthik chepudira
## Description: combiner code for HW3.2 Part 3
import sys
prev_key=-1
count=0
total_count =0 #keep count of the keys processed
for line in sys.stdin:
    key,value=line.split("\t",1)
    if key == "*":
        total_count+= eval(value) #increment the total key count from the mapper output.
    else:
        if prev_key != key:
            if prev_key !=-1:
                print prev_key+"\t"+str(count) #print key,value pair
                count = 0 #reset count
                prev_key = key # reset key
                count+=eval(value)
            else:
                count+=eval(value) # increment count
print prev_key+"\t"+str(count) # print the last key
print "*"+"\t"+str(total_count) #print the total key count.
sys.stderr.write("reporter:counter:MRCOUNTER,Combiner,1\n") # increment the counter for Combiner
```

Overwriting combiner3.2_3.py

```
In [11]: !chmod a+x combiner3.2_3.py
```

```
In [12]: #remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout3
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2_2" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=2 \
-input /tmp/3.1/Consumer_Complaints.csv \
-output /tmp/3.2/streamout3 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_2.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_2.py' \
-combiner '/usr/bin/python2.6 /usr/tmp/ds261/HW3/combiner3.2_3.py'
```

16/02/01 03:48:35 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Empty interval = 0 minutes

```
tes.
Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/streamout3' to
trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Curr
ent
WARNING: Use "yarn jar" to launch YARN applications.
packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-st
reaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob4395809441324277872.j
ar tmpDir=null
16/02/01 03:48:41 INFO impl.TimelineClientImpl: Timeline service ad
dress: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/01 03:48:42 INFO client.RMProxy: Connecting to ResourceManage
r at sandbox.hortonworks.com/192.168.214.133:8050
16/02/01 03:48:43 INFO impl.TimelineClientImpl: Timeline service ad
dress: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/01 03:48:43 INFO client.RMProxy: Connecting to ResourceManage
r at sandbox.hortonworks.com/192.168.214.133:8050
16/02/01 03:48:44 INFO mapred.FileInputFormat: Total input paths to
process : 1
16/02/01 03:48:44 INFO mapreduce.JobSubmitter: number of splits:2
16/02/01 03:48:44 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454250779452_0041
16/02/01 03:48:45 INFO impl.YarnClientImpl: Submitted application a
pplication_1454250779452_0041
16/02/01 03:48:45 INFO mapreduce.Job: The url to track the job: htt
p://sandbox.hortonworks.com:8088/proxy/application_1454250779452_00
41/
16/02/01 03:48:45 INFO mapreduce.Job: Running job: job_145425077945
2_0041
16/02/01 03:48:56 INFO mapreduce.Job: Job job_1454250779452_0041 ru
nning in uber mode : false
16/02/01 03:48:56 INFO mapreduce.Job: map 0% reduce 0%
16/02/01 03:49:13 INFO mapreduce.Job: map 2% reduce 0%
16/02/01 03:49:16 INFO mapreduce.Job: map 3% reduce 0%
16/02/01 03:49:19 INFO mapreduce.Job: map 4% reduce 0%
16/02/01 03:49:22 INFO mapreduce.Job: map 5% reduce 0%
16/02/01 03:49:25 INFO mapreduce.Job: map 6% reduce 0%
16/02/01 03:49:28 INFO mapreduce.Job: map 7% reduce 0%
16/02/01 03:49:31 INFO mapreduce.Job: map 8% reduce 0%
16/02/01 03:49:34 INFO mapreduce.Job: map 9% reduce 0%
16/02/01 03:49:37 INFO mapreduce.Job: map 10% reduce 0%
16/02/01 03:49:40 INFO mapreduce.Job: map 11% reduce 0%
16/02/01 03:49:43 INFO mapreduce.Job: map 12% reduce 0%
16/02/01 03:49:47 INFO mapreduce.Job: map 13% reduce 0%
16/02/01 03:49:50 INFO mapreduce.Job: map 14% reduce 0%
16/02/01 03:49:53 INFO mapreduce.Job: map 15% reduce 0%
16/02/01 03:49:56 INFO mapreduce.Job: map 16% reduce 0%
16/02/01 03:49:59 INFO mapreduce.Job: map 18% reduce 0%
16/02/01 03:50:02 INFO mapreduce.Job: map 19% reduce 0%
16/02/01 03:50:05 INFO mapreduce.Job: map 20% reduce 0%
16/02/01 03:50:08 INFO mapreduce.Job: map 21% reduce 0%
```

```
16/02/01 03:50:11 INFO mapreduce.Job: map 22% reduce 0%
16/02/01 03:50:14 INFO mapreduce.Job: map 24% reduce 0%
16/02/01 03:50:17 INFO mapreduce.Job: map 25% reduce 0%
16/02/01 03:50:20 INFO mapreduce.Job: map 26% reduce 0%
16/02/01 03:50:23 INFO mapreduce.Job: map 27% reduce 0%
16/02/01 03:50:26 INFO mapreduce.Job: map 28% reduce 0%
16/02/01 03:50:27 INFO mapreduce.Job: map 29% reduce 0%
16/02/01 03:50:30 INFO mapreduce.Job: map 30% reduce 0%
16/02/01 03:50:33 INFO mapreduce.Job: map 31% reduce 0%
16/02/01 03:50:35 INFO mapreduce.Job: map 32% reduce 0%
16/02/01 03:50:39 INFO mapreduce.Job: map 33% reduce 0%
16/02/01 03:50:42 INFO mapreduce.Job: map 34% reduce 0%
16/02/01 03:50:45 INFO mapreduce.Job: map 35% reduce 0%
16/02/01 03:50:47 INFO mapreduce.Job: map 36% reduce 0%
16/02/01 03:50:50 INFO mapreduce.Job: map 37% reduce 0%
16/02/01 03:50:53 INFO mapreduce.Job: map 38% reduce 0%
16/02/01 03:50:58 INFO mapreduce.Job: map 39% reduce 0%
16/02/01 03:51:01 INFO mapreduce.Job: map 41% reduce 0%
16/02/01 03:51:04 INFO mapreduce.Job: map 42% reduce 0%
16/02/01 03:51:07 INFO mapreduce.Job: map 43% reduce 0%
16/02/01 03:51:10 INFO mapreduce.Job: map 44% reduce 0%
16/02/01 03:51:13 INFO mapreduce.Job: map 45% reduce 0%
16/02/01 03:51:16 INFO mapreduce.Job: map 46% reduce 0%
16/02/01 03:51:19 INFO mapreduce.Job: map 47% reduce 0%
16/02/01 03:51:22 INFO mapreduce.Job: map 48% reduce 0%
16/02/01 03:51:25 INFO mapreduce.Job: map 49% reduce 0%
16/02/01 03:51:28 INFO mapreduce.Job: map 50% reduce 0%
16/02/01 03:51:31 INFO mapreduce.Job: map 51% reduce 0%
16/02/01 03:51:34 INFO mapreduce.Job: map 52% reduce 0%
16/02/01 03:51:37 INFO mapreduce.Job: map 53% reduce 0%
16/02/01 03:51:40 INFO mapreduce.Job: map 54% reduce 0%
16/02/01 03:51:43 INFO mapreduce.Job: map 55% reduce 0%
16/02/01 03:51:46 INFO mapreduce.Job: map 56% reduce 0%
16/02/01 03:51:49 INFO mapreduce.Job: map 57% reduce 0%
16/02/01 03:51:52 INFO mapreduce.Job: map 59% reduce 0%
16/02/01 03:51:55 INFO mapreduce.Job: map 60% reduce 0%
16/02/01 03:51:58 INFO mapreduce.Job: map 61% reduce 0%
16/02/01 03:52:01 INFO mapreduce.Job: map 62% reduce 0%
16/02/01 03:52:05 INFO mapreduce.Job: map 63% reduce 0%
16/02/01 03:52:08 INFO mapreduce.Job: map 64% reduce 0%
16/02/01 03:52:12 INFO mapreduce.Job: map 81% reduce 0%
16/02/01 03:52:14 INFO mapreduce.Job: map 82% reduce 0%
16/02/01 03:52:23 INFO mapreduce.Job: map 83% reduce 0%
16/02/01 03:52:31 INFO mapreduce.Job: map 83% reduce 8%
16/02/01 03:52:32 INFO mapreduce.Job: map 83% reduce 17%
16/02/01 03:52:34 INFO mapreduce.Job: map 100% reduce 17%
16/02/01 03:52:37 INFO mapreduce.Job: map 100% reduce 100%
16/02/01 03:52:38 INFO mapreduce.Job: Job job_1454250779452_0041 co
mpleted successfully
16/02/01 03:52:38 INFO mapreduce.Job: Counters: 52
```

File System Counters

FILE: Number of bytes read=4284
 FILE: Number of bytes written=520662
 FILE: Number of read operations=0
 FILE: Number of large read operations=0
 FILE: Number of write operations=0
 HDFS: Number of bytes read=51012846
 HDFS: Number of bytes written=2032
 HDFS: Number of read operations=12
 HDFS: Number of large read operations=0
 HDFS: Number of write operations=4

Job Counters

Launched map tasks=2
 Launched reduce tasks=2
 Data-local map tasks=2
 Total time spent by all maps in occupied slots (ms)
 =410332
 Total time spent by all reduces in occupied slots (ms)=45334
 Total time spent by all map tasks (ms)=410332
 Total time spent by all reduce tasks (ms)=45334
 Total vcore-seconds taken by all map tasks=410332
 Total vcore-seconds taken by all reduce tasks=45334
 Total megabyte-seconds taken by all map tasks=10258
 3000
 Total megabyte-seconds taken by all reduce tasks=11
 333500

Map-Reduce Framework

Map input records=312919
 Map output records=1002221
 Map output bytes=11872708
 Map output materialized bytes=4296
 Input split bytes=238
 Combine input records=1002221
 Combine output records=282
 Reduce input groups=157
 Reduce shuffle bytes=4296
 Reduce input records=282
 Reduce output records=153
 Spilled Records=564
 Shuffled Maps =4
 Failed Shuffles=0
 Merged Map outputs=4
 GC time elapsed (ms)=1068
 CPU time spent (ms)=380050
 Physical memory (bytes) snapshot=689311744
 Virtual memory (bytes) snapshot=4096471040
 Total committed heap usage (bytes)=495452160

MRCounter

Combiner=4

```
Mapper=2
Reducer=2
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=51012608
File Output Format Counters
  Bytes Written=2032
16/02/01 03:52:38 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout3
```

3.2 Part 3 Response

The value of the user defined Mapper Counter is: 2

The value of the user defined Combiner Counter is: 4

The value of the user defined Reducer Counter is: 2

```
MRCounter
  Combiner=4
  Mapper=2
  Reducer=2
```

3.2 Part 4

Using a single reducer: What are the top 50 most frequent terms in your word count analysis? Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items).

mapper 3.2_4

```
In [13]: %%writefile mapper3.2_4.py
#!/usr/bin/env python
## mapper3.2_4.py
## Author: karthik chepudira
## Description: mapper code for HW3.2 part 4
import sys
import re

for line in sys.stdin:
    if not line in ['\n', '\r\n']: #check for blank lines
        print line.rstrip() #print line

sys.stderr.write("reporter:counter:MRCOUNTER,Mapper,1\n") # increment the counter for Mapper
```

Overwriting mapper3.2_4.py

Reducer 3.2_4

```
In [7]: %%writefile reducer3.2_4.py
#!/usr/bin/env python
## reducer3.2_4.py
## Author: karthik chepudira
## Description: reducer code for HW3.2 part 4
import sys
import Queue

n_max, n_min = 50, 10 #number of most and least frequent items
q_min = Queue.Queue(n_min) # implement a queue for least frequent items
a_max = [] # list of most frequent items
total_count=0

for line in sys.stdin:
    word,count=line.rstrip().split("\t",1)
    if word == "*":
        total_count += eval(count) # Total number of records
    else:
        # put the biggest
        if len(a_max) < n_max:
            a_max.append(word+","+count) #append the most frequent items to the list
        if q_min.full():
            q_min.get()
            q_min.put(word+","+count) #insert next least frequent item to queue

print '\n%d Most Frequent records:' %n_max
for record in a_max:
    word,count=record.split(",",1)
    print record+","+str((1.0*eval(count))/(total_count)) # calculate relative frequency and print record

print '\n%d Least Frequent records:' %n_min
for i in range(n_min):
    word,count=q_min.get().split(",",1)
    print word+","+count+","+str((1.0*eval(count))/(total_count)) # calculate relative frequency and print record

sys.stderr.write("reporter:counter:MRCOUNTER,Reducer,1\n") #increment the Reducer Counter
```

Overwriting reducer3.2_4.py

```
In [8]: !chmod a+x mapper3.2_4.py reducer3.2_4.py
```


MapReduce Job 1 to calculate frequency for all items and Total Count

```
In [21]: ## remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout4
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2_4" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=1 \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keypartitioner.options=-k1,1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.text.key.comparator.options="-k1 -k2nr" \
-input /tmp/3.1/Consumer_Complaints.csv \
-output /tmp/3.2/streamout4 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_2.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_2.py' \
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner \
-combiner '/usr/bin/python2.6 /usr/tmp/ds261/HW3/combiner3.2_3.py'
```

16/02/01 04:32:50 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.

Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/streamout4' to trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Current

WARNING: Use "yarn jar" to launch YARN applications.

packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob85037138404013378.jar tmpDir=null

16/02/01 04:32:56 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 04:32:56 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/01 04:32:57 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 04:32:57 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/01 04:32:58 INFO mapred.FileInputFormat: Total input paths to process : 1

16/02/01 04:32:58 INFO mapreduce.JobSubmitter: number of splits:2

16/02/01 04:32:59 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454250779452_0044

16/02/01 04:33:00 INFO impl.YarnClientImpl: Submitted application a

pplication_1454250779452_0044

16/02/01 04:33:00 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454250779452_0044/

16/02/01 04:33:00 INFO mapreduce.Job: Running job: job_1454250779452_0044

16/02/01 04:33:14 INFO mapreduce.Job: Job job_1454250779452_0044 running in uber mode : false

16/02/01 04:33:14 INFO mapreduce.Job: map 0% reduce 0%

16/02/01 04:33:30 INFO mapreduce.Job: map 1% reduce 0%

16/02/01 04:33:31 INFO mapreduce.Job: map 2% reduce 0%

16/02/01 04:33:33 INFO mapreduce.Job: map 3% reduce 0%

16/02/01 04:33:36 INFO mapreduce.Job: map 4% reduce 0%

16/02/01 04:33:37 INFO mapreduce.Job: map 5% reduce 0%

16/02/01 04:33:40 INFO mapreduce.Job: map 6% reduce 0%

16/02/01 04:33:43 INFO mapreduce.Job: map 7% reduce 0%

16/02/01 04:33:45 INFO mapreduce.Job: map 8% reduce 0%

16/02/01 04:33:48 INFO mapreduce.Job: map 9% reduce 0%

16/02/01 04:33:51 INFO mapreduce.Job: map 10% reduce 0%

16/02/01 04:33:52 INFO mapreduce.Job: map 11% reduce 0%

16/02/01 04:33:55 INFO mapreduce.Job: map 12% reduce 0%

16/02/01 04:33:57 INFO mapreduce.Job: map 13% reduce 0%

16/02/01 04:34:01 INFO mapreduce.Job: map 14% reduce 0%

16/02/01 04:34:04 INFO mapreduce.Job: map 16% reduce 0%

16/02/01 04:34:07 INFO mapreduce.Job: map 17% reduce 0%

16/02/01 04:34:10 INFO mapreduce.Job: map 18% reduce 0%

16/02/01 04:34:13 INFO mapreduce.Job: map 19% reduce 0%

16/02/01 04:34:16 INFO mapreduce.Job: map 21% reduce 0%

16/02/01 04:34:20 INFO mapreduce.Job: map 22% reduce 0%

16/02/01 04:34:23 INFO mapreduce.Job: map 23% reduce 0%

16/02/01 04:34:26 INFO mapreduce.Job: map 24% reduce 0%

16/02/01 04:34:29 INFO mapreduce.Job: map 25% reduce 0%

16/02/01 04:34:32 INFO mapreduce.Job: map 26% reduce 0%

16/02/01 04:34:35 INFO mapreduce.Job: map 27% reduce 0%

16/02/01 04:34:38 INFO mapreduce.Job: map 28% reduce 0%

16/02/01 04:34:41 INFO mapreduce.Job: map 29% reduce 0%

16/02/01 04:34:44 INFO mapreduce.Job: map 30% reduce 0%

16/02/01 04:34:47 INFO mapreduce.Job: map 32% reduce 0%

16/02/01 04:34:50 INFO mapreduce.Job: map 33% reduce 0%

16/02/01 04:34:53 INFO mapreduce.Job: map 34% reduce 0%

16/02/01 04:34:56 INFO mapreduce.Job: map 36% reduce 0%

16/02/01 04:34:59 INFO mapreduce.Job: map 37% reduce 0%

16/02/01 04:35:02 INFO mapreduce.Job: map 38% reduce 0%

16/02/01 04:35:05 INFO mapreduce.Job: map 39% reduce 0%

16/02/01 04:35:08 INFO mapreduce.Job: map 40% reduce 0%

16/02/01 04:35:12 INFO mapreduce.Job: map 41% reduce 0%

16/02/01 04:35:18 INFO mapreduce.Job: map 42% reduce 0%

16/02/01 04:35:21 INFO mapreduce.Job: map 43% reduce 0%

16/02/01 04:35:27 INFO mapreduce.Job: map 44% reduce 0%

16/02/01 04:35:30 INFO mapreduce.Job: map 45% reduce 0%

```

16/02/01 04:35:33 INFO mapreduce.Job: map 46% reduce 0%
16/02/01 04:35:37 INFO mapreduce.Job: map 47% reduce 0%
16/02/01 04:35:40 INFO mapreduce.Job: map 48% reduce 0%
16/02/01 04:35:43 INFO mapreduce.Job: map 49% reduce 0%
16/02/01 04:35:46 INFO mapreduce.Job: map 50% reduce 0%
16/02/01 04:35:50 INFO mapreduce.Job: map 51% reduce 0%
16/02/01 04:35:56 INFO mapreduce.Job: map 52% reduce 0%
16/02/01 04:35:59 INFO mapreduce.Job: map 53% reduce 0%
16/02/01 04:36:02 INFO mapreduce.Job: map 54% reduce 0%
16/02/01 04:36:05 INFO mapreduce.Job: map 55% reduce 0%
16/02/01 04:36:08 INFO mapreduce.Job: map 56% reduce 0%
16/02/01 04:36:11 INFO mapreduce.Job: map 57% reduce 0%
16/02/01 04:36:14 INFO mapreduce.Job: map 58% reduce 0%
16/02/01 04:36:17 INFO mapreduce.Job: map 59% reduce 0%
16/02/01 04:36:20 INFO mapreduce.Job: map 60% reduce 0%
16/02/01 04:36:23 INFO mapreduce.Job: map 61% reduce 0%
16/02/01 04:36:26 INFO mapreduce.Job: map 62% reduce 0%
16/02/01 04:36:29 INFO mapreduce.Job: map 64% reduce 0%
16/02/01 04:36:32 INFO mapreduce.Job: map 65% reduce 0%
16/02/01 04:36:38 INFO mapreduce.Job: map 66% reduce 0%
16/02/01 04:36:42 INFO mapreduce.Job: map 83% reduce 0%
16/02/01 04:36:55 INFO mapreduce.Job: map 100% reduce 0%
16/02/01 04:36:57 INFO mapreduce.Job: map 100% reduce 100%
16/02/01 04:36:57 INFO mapreduce.Job: Job job_1454250779452_0044 co
mpleted successfully
16/02/01 04:36:57 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=4266
        FILE: Number of bytes written=395015
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=51012846
        HDFS: Number of bytes written=2028
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=425688
        Total time spent by all reduces in occupied slots (
ms)=11795
        Total time spent by all map tasks (ms)=425688
        Total time spent by all reduce tasks (ms)=11795
        Total vcore-seconds taken by all map tasks=425688
        Total vcore-seconds taken by all reduce tasks=11795
        Total megabyte-seconds taken by all map tasks=10642

```

```

2000
Total megabyte-seconds taken by all reduce tasks=29
48750
Map-Reduce Framework
  Map input records=312919
  Map output records=1002221
  Map output bytes=12874929
  Map output materialized bytes=4272
  Input split bytes=238
  Combine input records=1002221
  Combine output records=280
  Reduce input groups=154
  Reduce shuffle bytes=4272
  Reduce input records=280
  Reduce output records=152
  Spilled Records=560
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=1015
  CPU time spent (ms)=398290
  Physical memory (bytes) snapshot=545005568
  Virtual memory (bytes) snapshot=3052392448
  Total committed heap usage (bytes)=382730240
MRCounter
  Combiner=2
  Mapper=2
  Reducer=1
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=51012608
File Output Format Counters
  Bytes Written=2028
16/02/01 04:36:57 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout4

```

MapReduce Job 2 to sort and filter for Most Frequent and Least Frequent Items

```
In [22]: ## remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout4_2
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2_4" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=1 \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keypartitioner.options=-k1,1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.text.key.comparator.options="-k2nr -k1" \
-input /tmp/3.2/streamout4/* \
-output /tmp/3.2/streamout4_2 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_4.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_4.py' \
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
#-combiner '/usr/bin/python2.6 /usr/tmp/ds261/HW3/combiner3.2_3.py'
```

16/02/01 04:37:26 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.

Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/streamout4_2' to trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Current

WARNING: Use "yarn jar" to launch YARN applications.

packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob8371278169569604881.jar tmpDir=null

16/02/01 04:37:31 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 04:37:32 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/01 04:37:32 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 04:37:32 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.214.133:8050

16/02/01 04:37:33 INFO mapred.FileInputFormat: Total input paths to process : 1

16/02/01 04:37:33 INFO mapreduce.JobSubmitter: number of splits:2

16/02/01 04:37:34 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454250779452_0045

16/02/01 04:37:34 INFO impl.YarnClientImpl: Submitted application application_1454250779452_0045

16/02/01 04:37:34 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454250779452_0045/

16/02/01 04:37:34 INFO mapreduce.Job: Running job: job_145425077945

2_0045

16/02/01 04:37:44 INFO mapreduce.Job: Job job_1454250779452_0045 running in uber mode : false

16/02/01 04:37:44 INFO mapreduce.Job: map 0% reduce 0%

16/02/01 04:37:55 INFO mapreduce.Job: map 100% reduce 0%

16/02/01 04:38:03 INFO mapreduce.Job: map 100% reduce 100%

16/02/01 04:38:03 INFO mapreduce.Job: Job job_1454250779452_0045 completed successfully

16/02/01 04:38:03 INFO mapreduce.Job: Counters: 51

File System Counters

FILE: Number of bytes read=2490

FILE: Number of bytes written=390305

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=3276

HDFS: Number of bytes written=1963

HDFS: Number of read operations=9

HDFS: Number of large read operations=0

HDFS: Number of write operations=2

Job Counters

Launched map tasks=2

Launched reduce tasks=1

Data-local map tasks=2

Total time spent by all maps in occupied slots (ms)

=17428

Total time spent by all reduces in occupied slots (

ms)=5989

Total time spent by all map tasks (ms)=17428

Total time spent by all reduce tasks (ms)=5989

Total vcore-seconds taken by all map tasks=17428

Total vcore-seconds taken by all reduce tasks=5989

Total megabyte-seconds taken by all map tasks=43570

00

Total megabyte-seconds taken by all reduce tasks=14

97250

Map-Reduce Framework

Map input records=152

Map output records=152

Map output bytes=2180

Map output materialized bytes=2496

Input split bytes=234

Combine input records=0

Combine output records=0

Reduce input groups=152

Reduce shuffle bytes=2496

Reduce input records=152

Reduce output records=64

Spilled Records=304

Shuffled Maps =2

```

Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=285
CPU time spent (ms)=3710
Physical memory (bytes) snapshot=556621824
Virtual memory (bytes) snapshot=3073245184
Total committed heap usage (bytes)=375914496

MRCounter
  Mapper=2
  Reducer=1

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=3042

File Output Format Counters
  Bytes Written=1963

16/02/01 04:38:03 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout4_2

```

```
In [23]: !hdfs dfs -cat /tmp/3.2/streamout4_2/part-00000
```

```

50 Most Frequent records:
loan,119491,0.119226436537
modificationcollectionforeclosure,70487,0.0703309356538
account,57448,0.0573208051334
credit,50897,0.050784309617
payments,39993,0.039904452021
escrow,36767,0.0366855946654
servicing,36767,0.0366855946654
report,30546,0.0304783685003
incorrect,29069,0.0290046387067
information,29069,0.0290046387067
debt,27876,0.0278142801124
attempts,17972,0.0179322084295
collect,17972,0.0179322084295
contd,17972,0.0179322084295
owed,17972,0.0179322084295
closing,16205,0.0161691207211
management,16205,0.0161691207211
opening,16205,0.0161691207211
deposits,10555,0.0105316303123
withdrawals,10555,0.0105316303123
problems,9484,0.00946300159945
application,8868,0.00884836547701

```

communication,8671,0.00865180165213
tactics,8671,0.00865180165213
broker,8625,0.00860590350013
mortgage,8625,0.00860590350013
originator,8625,0.00860590350013
unable,8178,0.008159893197
billing,8158,0.00813993747873
disclosure,7655,0.00763805116447
verification,7655,0.00763805116447
disputes,6938,0.0069226386648
reporting,6560,0.00654547558967
lease,6337,0.00632296933105
caused,5663,0.00565046162565
funds,5663,0.00565046162565
low,5663,0.00565046162565
process,5505,0.00549281145139
managing,5006,0.00499491628077
improper,4966,0.00495500484425
companys,4858,0.00484724396564
investigation,4858,0.00484724396564
identity,4729,0.00471852958286
card,4407,0.00439724251885
get,4357,0.0043473532232
reportcredit,4357,0.0043473532232
score,4357,0.0043473532232
costs,4350,0.00434036872181
settlement,4350,0.00434036872181
interest,4238,0.00422861669954

10 Least Frequent records:

amount,98,9.77830194798e-05
credited,92,9.17963040014e-05
payment,92,9.17963040014e-05
checks,75,7.48339434794e-05
convenience,75,7.48339434794e-05
amt,71,7.08427998272e-05
day,71,7.08427998272e-05
disclosures,64,6.38582984358e-05
incorrectmissing,64,6.38582984358e-05
citibank,1,9.97785913059e-07

3.2 Part 4 Response

We used 2 Mapreduce jobs:

Job 1. First one to get frequencies of all the terms and the count of documents

Job 2. To sort the terms, terms frequencies by value, key columns
filter, print 50 most frequent terms and 10 least frequent terms

3.2.1 OPTIONAL

Using 2 reducers: What are the top 50 most frequent terms in your word count analysis? Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items).

MapReduce Job 1 to calculate frequency for all items and Total Count

```
In [1]: ## remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout5
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2.1" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=2 \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keypartitioner.options=-k1,1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapreduce.partition.keycomparator.options="-k1,1 -k2,2nr" \
-D mapred.text.key.comparator.options="-k1,1 -k2,2nr" \
-input /tmp/3.1/Consumer_Complaints.csv \
-output /tmp/3.2/streamout5 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_2.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_2.py' \
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner \
-combiner '/usr/bin/python2.6 /usr/tmp/ds261/HW3/combiner3.2_3.py'

rm -f /tmp/3.2/streamout5: No such file or directory
```

```
WARNING: Use "yarn jar" to launch YARN applications.
packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob2879157563458909815.jar tmpDir=null
16/02/01 21:30:02 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/01 21:30:02 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050
16/02/01 21:30:02 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/
16/02/01 21:30:02 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050
16/02/01 21:30:03 INFO mapred.FileInputFormat: Total input paths to process : 1
16/02/01 21:30:03 INFO mapreduce.JobSubmitter: number of splits:2
16/02/01 21:30:03 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454357458596_0001
16/02/01 21:30:04 INFO impl.YarnClientImpl: Submitted application application_1454357458596_0001
16/02/01 21:30:04 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454357458596_0001/
16/02/01 21:30:04 INFO mapreduce.Job: Running job: job_1454357458596_0001
16/02/01 21:30:15 INFO mapreduce.Job: Job job_1454357458596_0001 running in uber mode : false
16/02/01 21:30:15 INFO mapreduce.Job: map 0% reduce 0%
16/02/01 21:30:27 INFO mapreduce.Job: map 4% reduce 0%
16/02/01 21:30:30 INFO mapreduce.Job: map 6% reduce 0%
16/02/01 21:30:33 INFO mapreduce.Job: map 8% reduce 0%
16/02/01 21:30:36 INFO mapreduce.Job: map 10% reduce 0%
16/02/01 21:30:39 INFO mapreduce.Job: map 12% reduce 0%
16/02/01 21:30:42 INFO mapreduce.Job: map 15% reduce 0%
16/02/01 21:30:45 INFO mapreduce.Job: map 17% reduce 0%
16/02/01 21:30:48 INFO mapreduce.Job: map 19% reduce 0%
16/02/01 21:30:51 INFO mapreduce.Job: map 21% reduce 0%
16/02/01 21:30:54 INFO mapreduce.Job: map 24% reduce 0%
16/02/01 21:30:57 INFO mapreduce.Job: map 26% reduce 0%
16/02/01 21:31:00 INFO mapreduce.Job: map 28% reduce 0%
16/02/01 21:31:03 INFO mapreduce.Job: map 30% reduce 0%
16/02/01 21:31:06 INFO mapreduce.Job: map 32% reduce 0%
16/02/01 21:31:09 INFO mapreduce.Job: map 34% reduce 0%
16/02/01 21:31:12 INFO mapreduce.Job: map 35% reduce 0%
16/02/01 21:31:15 INFO mapreduce.Job: map 37% reduce 0%
16/02/01 21:31:18 INFO mapreduce.Job: map 39% reduce 0%
16/02/01 21:31:21 INFO mapreduce.Job: map 41% reduce 0%
16/02/01 21:31:24 INFO mapreduce.Job: map 43% reduce 0%
16/02/01 21:31:27 INFO mapreduce.Job: map 45% reduce 0%
16/02/01 21:31:30 INFO mapreduce.Job: map 48% reduce 0%
16/02/01 21:31:33 INFO mapreduce.Job: map 50% reduce 0%
```

```

16/02/01 21:31:36 INFO mapreduce.Job: map 52% reduce 0%
16/02/01 21:31:39 INFO mapreduce.Job: map 54% reduce 0%
16/02/01 21:31:42 INFO mapreduce.Job: map 57% reduce 0%
16/02/01 21:31:45 INFO mapreduce.Job: map 59% reduce 0%
16/02/01 21:31:48 INFO mapreduce.Job: map 61% reduce 0%
16/02/01 21:31:51 INFO mapreduce.Job: map 63% reduce 0%
16/02/01 21:31:54 INFO mapreduce.Job: map 64% reduce 0%
16/02/01 21:31:56 INFO mapreduce.Job: map 81% reduce 0%
16/02/01 21:31:57 INFO mapreduce.Job: map 82% reduce 0%
16/02/01 21:32:00 INFO mapreduce.Job: map 83% reduce 0%
16/02/01 21:32:11 INFO mapreduce.Job: map 100% reduce 8%
16/02/01 21:32:12 INFO mapreduce.Job: map 100% reduce 58%
16/02/01 21:32:13 INFO mapreduce.Job: map 100% reduce 100%
16/02/01 21:32:14 INFO mapreduce.Job: Job job_1454357458596_0001 completed successfully
16/02/01 21:32:14 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=4284
        FILE: Number of bytes written=523894
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=51012846
        HDFS: Number of bytes written=2032
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=214521
        Total time spent by all reduces in occupied slots (
ms)=24570
        Total time spent by all map tasks (ms)=214521
        Total time spent by all reduce tasks (ms)=24570
        Total vcore-seconds taken by all map tasks=214521
        Total vcore-seconds taken by all reduce tasks=24570
        Total megabyte-seconds taken by all map tasks=53630
250
        Total megabyte-seconds taken by all reduce tasks=61
42500
    Map-Reduce Framework
        Map input records=312919
        Map output records=1002221
        Map output bytes=12874929
        Map output materialized bytes=4296
        Input split bytes=238
        Combine input records=1002221

```

```
Combine output records=282
Reduce input groups=157
Reduce shuffle bytes=4296
Reduce input records=282
Reduce output records=153
Spilled Records=564
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=408
CPU time spent (ms)=197350
Physical memory (bytes) snapshot=688414720
Virtual memory (bytes) snapshot=4103147520
Total committed heap usage (bytes)=504889344

MRCounter
  Combiner=4
  Mapper=2
  Reducer=2

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=51012608
File Output Format Counters
  Bytes Written=2032
16/02/01 21:32:14 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout5
```

MapReduce Job 2 to sort and filter for Most Frequent and Least Frequent Items

```
In [9]: ## remove directory
!hdfs dfs -rm -r /tmp/3.2/streamout5_2
#run hadoop job
!hadoop jar /usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar \
-D mapred.job.name="HW3.2.1_2" \
-D mapred.map.tasks=2 \
-D mapred.reduce.tasks=1 \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keypartitioner.options=-k1,1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyFieldBasedComparator \
-D mapred.text.key.comparator.options="-k2nr -k1" \
-input /tmp/3.2/streamout5/* \
-output /tmp/3.2/streamout5_2 \
-mapper '/usr/bin/python2.6 /usr/tmp/ds261/HW3/mapper3.2_4.py' \
-reducer '/usr/bin/python2.6 /usr/tmp/ds261/HW3/reducer3.2_4.py' \
-partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
#-combiner '/usr/bin/python2.6 /usr/tmp/ds261/HW3/combiner3.2_3.py'
```

16/02/01 21:39:29 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 360 minutes, Emptier interval = 0 minutes.

Moved: 'hdfs://sandbox.hortonworks.com:8020/tmp/3.2/streamout5_2' to trash at: hdfs://sandbox.hortonworks.com:8020/user/root/.Trash/Current

WARNING: Use "yarn jar" to launch YARN applications.

packageJobJar: [] [/usr/hdp/2.3.0.0-2130/hadoop-mapreduce/hadoop-streaming-2.7.1.2.3.0.0-2130.jar] /tmp/streamjob2876634080924659687.jar tmpDir=null

16/02/01 21:39:32 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 21:39:32 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050

16/02/01 21:39:33 INFO impl.TimelineClientImpl: Timeline service address: http://sandbox.hortonworks.com:8188/ws/v1/timeline/

16/02/01 21:39:33 INFO client.RMProxy: Connecting to Resource Manager at sandbox.hortonworks.com/192.168.177.132:8050

16/02/01 21:39:33 INFO mapred.FileInputFormat: Total input paths to process : 2

16/02/01 21:39:33 INFO mapreduce.JobSubmitter: number of splits:2

16/02/01 21:39:33 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454357458596_0003

16/02/01 21:39:34 INFO impl.YarnClientImpl: Submitted application application_1454357458596_0003

16/02/01 21:39:34 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1454357458596_0003/

16/02/01 21:39:34 INFO mapreduce.Job: Running job: job_145435745859

```

6_0003
16/02/01 21:39:39 INFO mapreduce.Job: Job job_1454357458596_0003 ru
nning in uber mode : false
16/02/01 21:39:39 INFO mapreduce.Job: map 0% reduce 0%
16/02/01 21:39:45 INFO mapreduce.Job: map 100% reduce 0%
16/02/01 21:39:51 INFO mapreduce.Job: map 100% reduce 100%
16/02/01 21:39:51 INFO mapreduce.Job: Job job_1454357458596_0003 co
mpleted successfully
16/02/01 21:39:51 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=2497
        FILE: Number of bytes written=390325
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=2266
        HDFS: Number of bytes written=1963
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=8388
        Total time spent by all reduces in occupied slots (
ms)=2962
        Total time spent by all map tasks (ms)=8388
        Total time spent by all reduce tasks (ms)=2962
        Total vcore-seconds taken by all map tasks=8388
        Total vcore-seconds taken by all reduce tasks=2962
        Total megabyte-seconds taken by all map tasks=20970
00
        Total megabyte-seconds taken by all reduce tasks=74
0500
    Map-Reduce Framework
        Map input records=153
        Map output records=153
        Map output bytes=2185
        Map output materialized bytes=2503
        Input split bytes=234
        Combine input records=0
        Combine output records=0
        Reduce input groups=153
        Reduce shuffle bytes=2503
        Reduce input records=153
        Reduce output records=64
        Spilled Records=306
        Shuffled Maps =2

```

```

Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=150
CPU time spent (ms)=2180
Physical memory (bytes) snapshot=558735360
Virtual memory (bytes) snapshot=3083964416
Total committed heap usage (bytes)=371195904

MRCounter
  Mapper=2
  Reducer=1

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=2032

File Output Format Counters
  Bytes Written=1963

16/02/01 21:39:51 INFO streaming.StreamJob: Output directory: /tmp/
3.2/streamout5_2

```

```
In [12]: !hdfs dfs -cat /tmp/3.2/streamout5_2/part*
```

```

50 Most Frequent records:
loan,119491,0.119226436537
modificationcollectionforeclosure,70487,0.0703309356538
account,57448,0.0573208051334
credit,50897,0.050784309617
payments,39993,0.039904452021
escrow,36767,0.0366855946654
servicing,36767,0.0366855946654
report,30546,0.0304783685003
incorrect,29069,0.0290046387067
information,29069,0.0290046387067
debt,27876,0.0278142801124
attempts,17972,0.0179322084295
collect,17972,0.0179322084295
contd,17972,0.0179322084295
owed,17972,0.0179322084295
closing,16205,0.0161691207211
management,16205,0.0161691207211
opening,16205,0.0161691207211
deposits,10555,0.0105316303123
withdrawals,10555,0.0105316303123
problems,9484,0.00946300159945
application,8868,0.00884836547701

```

communication,8671,0.00865180165213
 tactics,8671,0.00865180165213
 broker,8625,0.00860590350013
 mortgage,8625,0.00860590350013
 originator,8625,0.00860590350013
 unable,8178,0.008159893197
 billing,8158,0.00813993747873
 disclosure,7655,0.00763805116447
 verification,7655,0.00763805116447
 disputes,6938,0.0069226386648
 reporting,6560,0.00654547558967
 lease,6337,0.00632296933105
 caused,5663,0.00565046162565
 funds,5663,0.00565046162565
 low,5663,0.00565046162565
 process,5505,0.00549281145139
 managing,5006,0.00499491628077
 improper,4966,0.00495500484425
 companys,4858,0.00484724396564
 investigation,4858,0.00484724396564
 identity,4729,0.00471852958286
 card,4407,0.00439724251885
 get,4357,0.0043473532232
 reportcredit,4357,0.0043473532232
 score,4357,0.0043473532232
 costs,4350,0.00434036872181
 settlement,4350,0.00434036872181
 interest,4238,0.00422861669954

10 Least Frequent records:
 amount,98,9.77830194798e-05
 credited,92,9.17963040014e-05
 payment,92,9.17963040014e-05
 checks,75,7.48339434794e-05
 convenience,75,7.48339434794e-05
 amt,71,7.08427998272e-05
 day,71,7.08427998272e-05
 disclosures,64,6.38582984358e-05
 incorrectmissing,64,6.38582984358e-05
 citibank,1,9.97785913059e-07

HW3.3. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

For this homework use the online browsing behavior dataset located at:

```
https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0
```

Each line in this dataset represents a browsing session of a customer.

On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData

FRO11987 ELE17451 ELE89019 SNA90258 GRO99222

GRO99222 GRO12298 FRO12685 ELE91550 SNA11465 ELE26917 ELE52966 FRO90334 SNA30755
ELE17451 FRO84225 SNA80192

ELE17451 GRO73461 DAI22896 SNA99873 FRO86643

ELE17451 ELE37798 FRO86643 GRO56989 ELE23393 SNA11465

ELE17451 SNA69641 FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

Do some exploratory data analysis of this dataset.

How many unique items are available from this supplier?

Using a single reducer: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

3.3 EDA Using Python

```
In [39]: # Use Python to do EDA
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
%matplotlib inline
```

We'll use Python to perform some exploratory data analysis before we use the MapReduce framework.


```
In [181]: def hw_3_3_eda():
# Items dictionary
items = Counter()
# Basket size variable
basket = []
# Iterate over records
with open("ProductPurchaseData.txt", "r") as f:
    for line in f:
        # Split into individual items
        record = line.strip().split()
        # Add items to dictionary and to basket
        items += Counter(record)
        basket.append(len(record))
# Print result and plot basket size distribution
print ("The number of total items: {}".format(sum(items.values
()))))
print ("The number of unique items: {}".format(len(items)))
print ("The average size of the basket: {:.2f}".format(np.mean
(basket)))
plt.figure(figsize=(14,8))
sns.distplot(basket, bins=15)

hw_3_3_eda()
```

The number of total items: 380824.

The number of unique items: 12592.

The average size of the basket: 12.24.



3.3 MapReduce

Strategy:

- 1st stage:
Mapper: purchase_record -> (item, 1, basketsize) (including (*,count) for order inversion in relative frequency calculation)
Reducer: (item, count) -> (item, total count, relative frequency, largest basketsize)
- 2nd stage:
 Sort 1st stage output by total count.

```
In [196]: !hdfs dfs -put ProductPurchaseData.txt /user/konniam/week_03/
```

```
16/01/29 23:59:04 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

```
In [42]: %%writefile mapper_3_3.py
#!/usr/bin/env python
# Mapper for 3.3
# Author: Konniam Chan
# Shopping cart analysis
# Output goes to 1 reducer
import sys

# Count how many mappers called
sys.stderr.write("reporter:counter:MR-Counter,Num_mapper_calls,1\n"
)

for line in sys.stdin:
    record = line.strip().split()
    basketsize = len(record)
    # Assume 1 reducer in the next step
    # Emit count and basketsize, (*) stands for the total
    for item in record:
        print '%s\t%s\t%s' % ('*', 1, basketsize)
        print '%s\t%s\t%s' % (item, 1, basketsize)
```

Overwriting mapper_3_3.py

```
In [43]: %%writefile combiner_3_3.py
#!/usr/bin/env python
## Combiner for 3.3
# Author: Konniam Chan
# Shopping cart analysis
import sys
current_item = None
current_count = 0
current_basketsize = 0
item = None

# Count how many combiners called
sys.stderr.write("reporter:counter:MR-Counter,Num_combiner_calls,1\n")

for line in sys.stdin:
    # Obtain item and intermediate counts
    line = line.strip()
    item, count, basketsize = line.split('\t')
    count = int(count)
    basketsize = int(basketsize)
    # Rely on sorting to increment item counts
    if current_item == item:
        current_count += count
        current_basketsize = max(current_basketsize, basketsize)
    else:
        if current_item:
            print '%s\t%s\t%s' % (current_item, current_count, current_basketsize)
        current_count = count
        current_item = item
        current_basketsize = basketsize

# Output last item
if current_item == item:
    print '%s\t%s\t%s' % (current_item, current_count, current_basketsize)
```

Overwriting combiner_3_3.py

```
In [45]: %%writefile reducer_3_3.py
#!/usr/bin/env python
# Reducer for 3.3
# Author: Konniam Chan
# Shopping cart analysis
# Assume 1 reducer
from __future__ import division
import sys
current_item = None
current_count = 0
current_basketsize = 0
item = None

# Count how many reducers called
sys.stderr.write("reporter:counter:MR-Counter,Num_reducer_calls,1\n")

for line in sys.stdin:
    # Obtain item and intermediate counts
    line = line.strip()
    item, count, basketsize = line.split('\t')
    count = int(count)
    basketsize = int(basketsize)
    # Rely on sorting to increment item counts
    if current_item == item:
        current_count += count
        current_basketsize = max(current_basketsize, basketsize)
    else:
        # Obtain total count and save to memory
        if current_item == '*':
            total = current_count
        # Compute counts and relative frequencies for regular items
        elif current_item:
            # Output (item, count, relative frequency)
            print '%s\t%s\t%s\t%s' % (current_item, current_count,
current_count / total, current_basketsize)
            current_count = count
            current_item = item
            current_basketsize = basketsize

# Output last item
if current_item == item:
    print '%s\t%s\t%s\t%s' % (current_item, current_count, current_
count / total, current_basketsize)
```

Writing reducer_3_3.py

```
In [47]: !chmod a+x *_3_3.py
```

```
In [48]: # Shopping cart analysis, with counters for number of mappers, combiners, and reducers
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_3
# Use 4 mapper and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.jar \
-D mapreduce.job.maps=4 \
-D mapreduce.job.reduces=1 \
-mapper $PWD/mapper_3_3.py \
-reducer $PWD/reducer_3_3.py \
-combiner $PWD/combiner_3_3.py \
-input /user/konniam/week_03/ProductPurchaseData.txt \
-output /user/konniam/week_03/hw_3_3
```

16/01/31 17:30:35 WARN util.NativeCodeLoader: Unable to load native -hadoop library for your platform... using builtin-java classes where applicable

16/01/31 17:30:38 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.

Deleted /user/konniam/week_03/hw_3_3

16/01/31 17:30:41 WARN util.NativeCodeLoader: Unable to load native -hadoop library for your platform... using builtin-java classes where applicable

packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/hadoop-unjar5571571146201173397/] [] /var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/streamjob6445029810789433212.jar tmpDir=null

16/01/31 17:30:42 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

16/01/31 17:30:43 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

16/01/31 17:30:44 INFO mapred.FileInputFormat: Total input paths to process : 1

16/01/31 17:30:45 INFO mapreduce.JobSubmitter: number of splits:4

16/01/31 17:30:45 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454277284610_0013

16/01/31 17:30:45 INFO impl.YarnClientImpl: Submitted application application_1454277284610_0013

16/01/31 17:30:45 INFO mapreduce.Job: The url to track the job: http://Konniam-MacBook-Air.local:8088/proxy/application_1454277284610_0013/

16/01/31 17:30:45 INFO mapreduce.Job: Running job: job_1454277284610_0013

16/01/31 17:30:56 INFO mapreduce.Job: Job job_1454277284610_0013 running in uber mode : false

16/01/31 17:30:56 INFO mapreduce.Job: map 0% reduce 0%

16/01/31 17:31:16 INFO mapreduce.Job: map 67% reduce 0%

16/01/31 17:31:18 INFO mapreduce.Job: map 100% reduce 0%

16/01/31 17:31:26 INFO mapreduce.Job: map 100% reduce 100%

```

16/01/31 17:31:27 INFO mapreduce.Job: Job job_1454277284610_0013 co
mpleted successfully
16/01/31 17:31:27 INFO mapreduce.Job: Counters: 52
  File System Counters
    FILE: Number of bytes read=399805
    FILE: Number of bytes written=1389203
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=3469551
    HDFS: Number of bytes written=404717
    HDFS: Number of read operations=15
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=4
    Launched reduce tasks=1
    Data-local map tasks=4
    Total time spent by all maps in occupied slots (ms)
=81420
    Total time spent by all reduces in occupied slots (
ms)=5922
    Total time spent by all map tasks (ms)=81420
    Total time spent by all reduce tasks (ms)=5922
    Total vcore-seconds taken by all map tasks=81420
    Total vcore-seconds taken by all reduce tasks=5922
    Total megabyte-seconds taken by all map tasks=83374
080
    Total megabyte-seconds taken by all reduce tasks=60
64128
  Map-Reduce Framework
    Map input records=31101
    Map output records=761648
    Map output bytes=7845208
    Map output materialized bytes=399823
    Input split bytes=472
    Combine input records=761648
    Combine output records=24800
    Reduce input groups=12593
    Reduce shuffle bytes=399823
    Reduce input records=24800
    Reduce output records=12592
    Spilled Records=49600
    Shuffled Maps =4
    Failed Shuffles=0
    Merged Map outputs=4
    GC time elapsed (ms)=769
    CPU time spent (ms)=0
    Physical memory (bytes) snapshot=0
    Virtual memory (bytes) snapshot=0

```

```

        Total committed heap usage (bytes)=812122112
MR-Counter
    Num_combiner_calls=4
    Num_mapper_calls=4
    Num_reducer_calls=1
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=3469079
File Output Format Counters
    Bytes Written=404717
16/01/31 17:31:27 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3

```

Top 50 Purchased Items MR:

```

In [4]: %%writefile reducer_top50.py
        #!/usr/bin/env python
        # Reducer for 3.3
        # Author: Konniam Chan
        # Apriori Algorithm
        import sys
        N = 50
        counter = 0
        for line in sys.stdin:
            # Ouput N items only
            if counter >= N:
                continue
            print line.strip()
            counter += 1

```

Overwriting reducer_top50.py

```

In [3]: !chmod a+x reducer_top50.py

```



```
In [5]: # Sort results by count from 1st MR stage
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_3_sorted_count
# Use 4 mapper and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=4 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educelib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keycomparator.options='-k2,2nr -k1,1' \
-mapper /bin/cat \
-reducer $PWD/reducer_top50.py \
-input /user/konniam/week_03/hw_3_3/part* \
-output /user/konniam/week_03/hw_3_3_sorted_count
```

```
16/02/03 11:23:52 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 11:23:53 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_3_sorted_count
16/02/03 11:23:54 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h5l_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar2509895711049292365/] [] /var/folders/l8/h5l_59852qscq403
fs6q0xlh0000gn/T/streamjob8194126334403388460.jar tmpDir=null
16/02/03 11:23:56 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:23:57 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:23:58 INFO mapred.FileInputFormat: Total input paths to
process : 1
16/02/03 11:23:58 INFO mapreduce.JobSubmitter: number of splits:4
16/02/03 11:23:59 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0022
16/02/03 11:23:59 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0022
16/02/03 11:23:59 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0022/
16/02/03 11:23:59 INFO mapreduce.Job: Running job: job_145448753435
8_0022
16/02/03 11:24:08 INFO mapreduce.Job: Job job_1454487534358_0022 ru
nning in uber mode : false
16/02/03 11:24:08 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:24:29 INFO mapreduce.Job:  map 50% reduce 0%
```

```

16/02/03 11:24:30 INFO mapreduce.Job: map 100% reduce 0%
16/02/03 11:24:38 INFO mapreduce.Job: map 100% reduce 100%
16/02/03 11:24:39 INFO mapreduce.Job: Job job_1454487534358_0022 completed successfully
16/02/03 11:24:40 INFO mapreduce.Job: Counters: 49

```

File System Counters

```

FILE: Number of bytes read=429907
FILE: Number of bytes written=1449877
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=408395
HDFS: Number of bytes written=1689
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

```

Job Counters

```

Launched map tasks=4
Launched reduce tasks=1
Data-local map tasks=4
Total time spent by all maps in occupied slots (ms)
=74628
Total time spent by all reduces in occupied slots (
ms)=6591
Total time spent by all map tasks (ms)=74628
Total time spent by all reduce tasks (ms)=6591
Total vcore-seconds taken by all map tasks=74628
Total vcore-seconds taken by all reduce tasks=6591
Total megabyte-seconds taken by all map tasks=76419
072
Total megabyte-seconds taken by all reduce tasks=67
49184

```

Map-Reduce Framework

```

Map input records=12592
Map output records=12592
Map output bytes=404717
Map output materialized bytes=429925
Input split bytes=448
Combine input records=0
Combine output records=0
Reduce input groups=12592
Reduce shuffle bytes=429925
Reduce input records=12592
Reduce output records=50
Spilled Records=25184
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=884
CPU time spent (ms)=0

```

```

Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=798490624
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=407947
File Output Format Counters
  Bytes Written=1689
16/02/03 11:24:40 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3_sorted_count

```

Top 50 Purchased Items Results:

Result format:

(Item, Count, Relative Frequency, Largest Basket in Which Item was Included)

```
In [6]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3_sorted_count/part-*
```

```

16/02/03 11:25:44 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

```

DAI62779	6667	0.0175067747831	36
FRO40251	3881	0.010191059387	37
ELE17451	3875	0.0101753040775	37
GRO73461	3602	0.00945843749344	37
SNA80324	3044	0.00799319370628	36
ELE32164	2851	0.0074863979161	36
DAI75645	2736	0.00718442114993	37
SNA45677	2455	0.0064465474865	36
FRO31317	2330	0.0061183118711	37
DAI85309	2293	0.00602115412894	36
ELE26917	2292	0.00601852824402	33
FRO80039	2233	0.00586360103355	34
GRO21487	2115	0.00555374661261	37
SNA99873	2083	0.00546971829507	35
GRO59710	2004	0.00526227338613	37
GRO71621	1920	0.00504169905258	36
FRO85978	1918	0.00503644728273	37
GRO30386	1840	0.00483162825872	36
ELE74009	1816	0.00476860702057	35
GRO56726	1784	0.00468457870302	36
DAI63921	1773	0.00465569396887	35

GRO46854	1756	0.00461105392517	37
ELE66600	1713	0.00449814087347	37
DAI83733	1712	0.00449551498855	36
FRO32293	1702	0.00446925613932	36
ELE66810	1697	0.0044561267147	36
SNA55762	1646	0.00432220658362	37
DAI22177	1627	0.00427231477008	36
FRO78087	1531	0.00402022981745	37
ELE99737	1516	0.0039808415436	33
ELE34057	1489	0.00390994265067	35
GRO94758	1489	0.00390994265067	36
FRO35904	1436	0.00377077074974	36
FRO53271	1420	0.00372875659097	36
SNA93860	1407	0.00369462008697	36
SNA90094	1390	0.00364998004327	36
GRO38814	1352	0.00355019641619	37
ELE56788	1345	0.00353181522173	36
GRO61133	1321	0.00346879398357	37
DAI88807	1316	0.00345566455896	35
ELE74482	1316	0.00345566455896	35
ELE59935	1311	0.00344253513434	36
SNA96271	1295	0.00340052097557	36
DAI43223	1290	0.00338739155095	37
ELE91337	1289	0.00338476566603	33
GRO15017	1275	0.0033480032771	37
DAI31081	1261	0.00331124088818	35
GRO81087	1220	0.00320357960633	36
DAI22896	1219	0.0032009537214	36
GRO85051	1214	0.00318782429679	34

Largest Basket MR

```
In [7]: %%writefile reducer_top1.py
#!/usr/bin/env python
# Reducer for 3.3
# Author: Konniam Chan
# Apriori Algorithm
import sys
N = 1
counter = 0
for line in sys.stdin:
    # Ouput N items only
    if counter >= N:
        continue
    print line.strip()
    counter += 1
```

Writing reducer_top1.py

In [8]: `!chmod a+x reducer_top1.py`

In [9]: `# Sort results by basketsize from 1st MR stage
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_3_sorted_basket
Use 4 mapper and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.jar \n
-D mapreduce.job.maps=4 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapreduce.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=4 \
-D mapreduce.partition.keycomparator.options='-k4,4nr -k1,1' \
-mapper /bin/cat \
-reducer $PWD/reducer_top1.py \
-input /user/konniam/week_03/hw_3_3/part* \
-output /user/konniam/week_03/hw_3_3_sorted_basket`

16/02/03 11:27:45 WARN util.NativeCodeLoader: Unable to load native -hadoop library for your platform... using builtin-java classes where applicable

16/02/03 11:27:46 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.

Deleted /user/konniam/week_03/hw_3_3_sorted_basket

16/02/03 11:27:48 WARN util.NativeCodeLoader: Unable to load native -hadoop library for your platform... using builtin-java classes where applicable

packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/hadoop-unjar6127660119398537226/] [] /var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/streamjob2025001027078091936.jar tmpDir=null

16/02/03 11:27:49 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

16/02/03 11:27:50 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032

16/02/03 11:27:51 INFO mapred.FileInputFormat: Total input paths to process : 1

16/02/03 11:27:51 INFO mapreduce.JobSubmitter: number of splits:4

16/02/03 11:27:51 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454487534358_0023

16/02/03 11:27:52 INFO impl.YarnClientImpl: Submitted application application_1454487534358_0023

16/02/03 11:27:52 INFO mapreduce.Job: The url to track the job: http://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358_0023/

16/02/03 11:27:52 INFO mapreduce.Job: Running job: job_1454487534358_0023

16/02/03 11:28:00 INFO mapreduce.Job: Job job_1454487534358_0023 running in uber mode : false

```

16/02/03 11:28:00 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:28:19 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 11:28:28 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 11:28:29 INFO mapreduce.Job: Job job_1454487534358_0023 co
mpleted successfully
16/02/03 11:28:29 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=442499
        FILE: Number of bytes written=1475061
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=408395
        HDFS: Number of bytes written=34
        HDFS: Number of read operations=15
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=4
        Launched reduce tasks=1
        Data-local map tasks=4
        Total time spent by all maps in occupied slots (ms)
=70134
        Total time spent by all reduces in occupied slots (
ms)=5322
        Total time spent by all map tasks (ms)=70134
        Total time spent by all reduce tasks (ms)=5322
        Total vcore-seconds taken by all map tasks=70134
        Total vcore-seconds taken by all reduce tasks=5322
        Total megabyte-seconds taken by all map tasks=71817
216
        Total megabyte-seconds taken by all reduce tasks=54
49728
    Map-Reduce Framework
        Map input records=12592
        Map output records=12592
        Map output bytes=417309
        Map output materialized bytes=442517
        Input split bytes=448
        Combine input records=0
        Combine output records=0
        Reduce input groups=12592
        Reduce shuffle bytes=442517
        Reduce input records=12592
        Reduce output records=1
        Spilled Records=25184
        Shuffled Maps =4
        Failed Shuffles=0
        Merged Map outputs=4
        GC time elapsed (ms)=860

```

```

CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=836763648
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=407947
File Output Format Counters
  Bytes Written=34
16/02/03 11:28:29 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3_sorted_basket

```

Largest Basket: 37

Note: the largest basket of 37 occurred for many items, and not specific to the one item shown (it is simply the one that came out in alphabetical order)

```
In [10]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3_sorted_basket/part-*
```

```

16/02/03 11:28:42 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
DAI35347          1060      0.00278343801861          37

```

Number of Unique Items: 12592

```
In [99]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3/part-* | wc -l
```

```

16/01/30 16:51:00 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
12592

```

3.3.1 OPTIONAL

Using 2 reducers: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

3.3.1 Strategy

The key is to direct total counts from the first mapper to the 2 reducers using two * keys ('*0' and '*1') for order inversion. The rest is same as above.

3.3.1 Mapper

```
In [1]: %%writefile mapper_3_3_1.py
#!/usr/bin/env python
# Mapper for 3.3.1
# Author: Konniam Chan
# Shopping cart analysis
# Output goes to 1 reducer
import sys

# Count how many mappers called
sys.stderr.write("reporter:counter:MR-Counter,Num_mapper_calls,1\n"
)

for line in sys.stdin:
    record = line.strip().split()
    basketsize = len(record)
    # Assume 2 reducers in the next step
    # Emit count and basketsize, (*) stands for the total
    for item in record:
        print '%s\t%s\t%s' % ('*0', 1, basketsize)
        print '%s\t%s\t%s' % ('*1', 1, basketsize)
        print '%s\t%s\t%s' % (item, 1, basketsize)
```

Writing mapper_3_3_1.py

3.3.1 Combiner

Can use the combiner from 3.3

3.3.1 Reducer


```
In [2]: %%writefile reducer_3_3_1.py
#!/usr/bin/env python
# Reducer for 3.3.1
# Author: Konniam Chan
# Shopping cart analysis
# Assume 2 reducers
from __future__ import division
import sys
current_item = None
current_count = 0
current_basketsize = 0
item = None

# Count how many reducers called
sys.stderr.write("reporter:counter:MR-Counter,Num_reducer_calls,1\n")

for line in sys.stdin:
    # Obtain item and intermediate counts
    line = line.strip()
    item, count, basketsize = line.split('\t')
    count = int(count)
    basketsize = int(basketsize)
    # Rely on sorting to increment item counts
    if current_item == item:
        current_count += count
        current_basketsize = max(current_basketsize, basketsize)
    else:
        # Obtain total count and save to memory
        if current_item in ['*0', '*1']:
            total = current_count
        # Compute counts and relative frequencies for regular items
        elif current_item:
            # Output (item, count, relative frequency)
            print '%s\t%s\t%s\t%s' % (current_item, current_count,
current_count / total, current_basketsize)
            current_count = count
            current_item = item
            current_basketsize = basketsize

# Output last item
if current_item == item:
    print '%s\t%s\t%s\t%s' % (current_item, current_count, current_
count / total, current_basketsize)

Writing reducer_3_3_1.py
```

```
In [3]: !chmod a+x *_3_3_1.py
```

3.3.1 1st Stage MR

```
In [4]: # Shopping cart analysis, with counters for number of mappers, comb
        iners, and reducers
        !hdfs dfs -rm -r /user/konniam/week_03/hw_3_3_1
        # Use 2 mappers and 2 reducers
        !hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
        jar \
        -D mapreduce.job.maps=2 \
        -D mapreduce.job.reduces=2 \
        -mapper $PWD/mapper_3_3_1.py \
        -reducer $PWD/reducer_3_3_1.py \
        -combiner $PWD/combiner_3_3.py \
        -input /user/konniam/week_03/ProductPurchaseData.txt \
        -output /user/konniam/week_03/hw_3_3_1
```

```
16/02/01 11:07:29 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

```
rm: `/user/konniam/week_03/hw_3_3_1': No such file or directory
```

```
16/02/01 11:07:32 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

```
packageJobJar: [/var/folders/l8/h5l_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar2188033139595451417/] [] /var/folders/l8/h5l_59852qscq403
fs6q0xlh0000gn/T/streamjob1897984548306338732.jar tmpDir=null
```

```
16/02/01 11:07:33 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
```

```
16/02/01 11:07:33 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
```

```
16/02/01 11:07:35 INFO mapred.FileInputFormat: Total input paths to
process : 1
```

```
16/02/01 11:07:35 INFO mapreduce.JobSubmitter: number of splits:2
```

```
16/02/01 11:07:35 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454277284610_0021
```

```
16/02/01 11:07:35 INFO impl.YarnClientImpl: Submitted application a
pplication_1454277284610_0021
```

```
16/02/01 11:07:36 INFO mapreduce.Job: The url to track the job: htt
p://Konniams-MacBook-Air.local:8088/proxy/application_1454277284610
_0021/
```

```
16/02/01 11:07:36 INFO mapreduce.Job: Running job: job_145427728461
0_0021
```

```
16/02/01 11:07:45 INFO mapreduce.Job: Job job_1454277284610_0021 ru
nning in uber mode : false
```

```
16/02/01 11:07:45 INFO mapreduce.Job: map 0% reduce 0%
```

```
16/02/01 11:07:59 INFO mapreduce.Job: map 67% reduce 0%
```

```
16/02/01 11:08:01 INFO mapreduce.Job: map 100% reduce 0%
```

```
16/02/01 11:08:10 INFO mapreduce.Job: map 100% reduce 50%
```

16/02/01 11:08:11 INFO mapreduce.Job: map 100% reduce 100%
 16/02/01 11:08:11 INFO mapreduce.Job: Job job_1454277284610_0021 completed successfully

16/02/01 11:08:12 INFO mapreduce.Job: Counters: 52

File System Counters

FILE: Number of bytes read=286952
 FILE: Number of bytes written=1045582
 FILE: Number of read operations=0
 FILE: Number of large read operations=0
 FILE: Number of write operations=0
 HDFS: Number of bytes read=3462103
 HDFS: Number of bytes written=404717
 HDFS: Number of read operations=12
 HDFS: Number of large read operations=0
 HDFS: Number of write operations=4

Job Counters

Launched map tasks=2
 Launched reduce tasks=2
 Data-local map tasks=2
 Total time spent by all maps in occupied slots (ms)
 =28778
 Total time spent by all reduces in occupied slots (ms)=12942
 Total time spent by all map tasks (ms)=28778
 Total time spent by all reduce tasks (ms)=12942
 Total vcore-seconds taken by all map tasks=28778
 Total vcore-seconds taken by all reduce tasks=12942
 Total megabyte-seconds taken by all map tasks=29468672
 Total megabyte-seconds taken by all reduce tasks=13252608

Map-Reduce Framework

Map input records=31101
 Map output records=1142472
 Map output bytes=11196576
 Map output materialized bytes=286964
 Input split bytes=236
 Combine input records=1142472
 Combine output records=17749
 Reduce input groups=12594
 Reduce shuffle bytes=286964
 Reduce input records=17749
 Reduce output records=12592
 Spilled Records=35498
 Shuffled Maps =4
 Failed Shuffles=0
 Merged Map outputs=4
 GC time elapsed (ms)=458
 CPU time spent (ms)=0
 Physical memory (bytes) snapshot=0

```

Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=595591168
MR-Counter
  Num_combiner_calls=4
  Num_mapper_calls=2
  Num_reducer_calls=2
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=3461867
File Output Format Counters
  Bytes Written=404717
16/02/01 11:08:12 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3_1

```

3.3.1 Top 50 Purchased Items MR:

```

In [1]: # Sort results by count from 1st MR stage
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_3_1_sorted_count
# Use 2 mappers and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educe.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=2 \
-D mapreduce.partition.keycomparator.options='-k2,2nr -k1,1' \
-mapper /bin/cat \
-reducer $PWD/reducer_top50.py \
-input /user/konniam/week_03/hw_3_3_1/part* \
-output /user/konniam/week_03/hw_3_3_1_sorted_count

```

```

16/02/03 11:42:55 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 11:42:56 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_3_1_sorted_count
16/02/03 11:42:58 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/ha

```

```

doop-unjar4222233125874937807/] [] /var/folders/l8/h5l_59852qscq403
fs6q0xlh0000gn/T/streamjob4453936316448288606.jar tmpDir=null
16/02/03 11:43:00 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:43:00 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:43:01 INFO mapred.FileInputFormat: Total input paths to
process : 2
16/02/03 11:43:01 INFO mapreduce.JobSubmitter: number of splits:2
16/02/03 11:43:02 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0024
16/02/03 11:43:02 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0024
16/02/03 11:43:02 INFO mapreduce.Job: The url to track the job: htt
p://Konniams-MacBook-Air.local:8088/proxy/application_1454487534358
_0024/
16/02/03 11:43:02 INFO mapreduce.Job: Running job: job_145448753435
8_0024
16/02/03 11:43:13 INFO mapreduce.Job: Job job_1454487534358_0024 ru
nning in uber mode : false
16/02/03 11:43:13 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:43:26 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 11:43:34 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 11:43:35 INFO mapreduce.Job: Job job_1454487534358_0024 co
mpleted successfully
16/02/03 11:43:35 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=429907
        FILE: Number of bytes written=1213837
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=404945
        HDFS: Number of bytes written=1689
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=20841
        Total time spent by all reduces in occupied slots (
ms)=4986
        Total time spent by all map tasks (ms)=20841
        Total time spent by all reduce tasks (ms)=4986
        Total vcore-seconds taken by all map tasks=20841
        Total vcore-seconds taken by all reduce tasks=4986
        Total megabyte-seconds taken by all map tasks=21341

```

184

Total megabyte-seconds taken by all reduce tasks=51

05664

Map-Reduce Framework

```

Map input records=12592
Map output records=12592
Map output bytes=404717
Map output materialized bytes=429913
Input split bytes=228
Combine input records=0
Combine output records=0
Reduce input groups=12592
Reduce shuffle bytes=429913
Reduce input records=12592
Reduce output records=50
Spilled Records=25184
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=462
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=497025024

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```

Bytes Read=404717

```

File Output Format Counters

```

Bytes Written=1689

```

```

16/02/03 11:43:35 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3_1_sorted_count

```

3.3.1 Top 50 Purchased Items Results:

Result format:

(Item, Count, Relative Frequency, Largest Basket in Which Item was Included)

```
In [2]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3_1_sorted_count/part-*
```

```

16/02/03 11:43:41 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

```

DAI62779	6667	0.0175067747831	36
FRO40251	3881	0.010191059387	37
ELE17451	3875	0.0101753040775	37
GRO73461	3602	0.00945843749344	37
SNA80324	3044	0.00799319370628	36
ELE32164	2851	0.0074863979161	36
DAI75645	2736	0.00718442114993	37
SNA45677	2455	0.0064465474865	36
FRO31317	2330	0.0061183118711	37
DAI85309	2293	0.00602115412894	36
ELE26917	2292	0.00601852824402	33
FRO80039	2233	0.00586360103355	34
GRO21487	2115	0.00555374661261	37
SNA99873	2083	0.00546971829507	35
GRO59710	2004	0.00526227338613	37
GRO71621	1920	0.00504169905258	36
FRO85978	1918	0.00503644728273	37
GRO30386	1840	0.00483162825872	36
ELE74009	1816	0.00476860702057	35
GRO56726	1784	0.00468457870302	36
DAI63921	1773	0.00465569396887	35
GRO46854	1756	0.00461105392517	37
ELE66600	1713	0.00449814087347	37
DAI83733	1712	0.00449551498855	36
FRO32293	1702	0.00446925613932	36
ELE66810	1697	0.0044561267147	36
SNA55762	1646	0.00432220658362	37
DAI22177	1627	0.00427231477008	36
FRO78087	1531	0.00402022981745	37
ELE99737	1516	0.0039808415436	33
ELE34057	1489	0.00390994265067	35
GRO94758	1489	0.00390994265067	36
FRO35904	1436	0.00377077074974	36
FRO53271	1420	0.00372875659097	36
SNA93860	1407	0.00369462008697	36
SNA90094	1390	0.00364998004327	36
GRO38814	1352	0.00355019641619	37
ELE56788	1345	0.00353181522173	36
GRO61133	1321	0.00346879398357	37
DAI88807	1316	0.00345566455896	35
ELE74482	1316	0.00345566455896	35
ELE59935	1311	0.00344253513434	36
SNA96271	1295	0.00340052097557	36
DAI43223	1290	0.00338739155095	37
ELE91337	1289	0.00338476566603	33
GRO15017	1275	0.0033480032771	37
DAI31081	1261	0.00331124088818	35
GRO81087	1220	0.00320357960633	36
DAI22896	1219	0.0032009537214	36
GRO85051	1214	0.00318782429679	34

3.3.1 Largest Basket MR

```
In [3]: # Sort results by basketsize from 1st MR stage
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_3_1_sorted_basket
# Use 2 mappers and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educe.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=4 \
-D mapreduce.partition.keycomparator.options='-k4,4nr -k1,1' \
-mapper /bin/cat \
-reducer $PWD/reducer_top1.py \
-input /user/konniam/week_03/hw_3_3_1/part* \
-output /user/konniam/week_03/hw_3_3_1_sorted_basket
```

```
16/02/03 11:43:56 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 11:43:57 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_3_1_sorted_basket
16/02/03 11:43:59 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar4460119617717925651/] [] /var/folders/l8/h51_59852qscq403
fs6q0xlh0000gn/T/streamjob3572732194774922216.jar tmpDir=null
16/02/03 11:44:00 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:44:00 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:44:01 INFO mapred.FileInputFormat: Total input paths to
process : 2
16/02/03 11:44:01 INFO mapreduce.JobSubmitter: number of splits:2
16/02/03 11:44:02 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0025
16/02/03 11:44:02 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0025
16/02/03 11:44:02 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0025/
16/02/03 11:44:02 INFO mapreduce.Job: Running job: job_145448753435
8_0025
16/02/03 11:44:11 INFO mapreduce.Job: Job job_1454487534358_0025 ru
nning in uber mode : false
```



```

16/02/03 11:44:11 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:44:23 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 11:44:30 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 11:44:31 INFO mapreduce.Job:  Job job_1454487534358_0025 co
mpleted successfully
16/02/03 11:44:31 INFO mapreduce.Job:  Counters: 49
    File System Counters
        FILE: Number of bytes read=442499
        FILE: Number of bytes written=1239021
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=404945
        HDFS: Number of bytes written=34
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=18463
        Total time spent by all reduces in occupied slots (
ms)=4643
        Total time spent by all map tasks (ms)=18463
        Total time spent by all reduce tasks (ms)=4643
        Total vcore-seconds taken by all map tasks=18463
        Total vcore-seconds taken by all reduce tasks=4643
        Total megabyte-seconds taken by all map tasks=18906
112
        Total megabyte-seconds taken by all reduce tasks=47
54432
    Map-Reduce Framework
        Map input records=12592
        Map output records=12592
        Map output bytes=417309
        Map output materialized bytes=442505
        Input split bytes=228
        Combine input records=0
        Combine output records=0
        Reduce input groups=12592
        Reduce shuffle bytes=442505
        Reduce input records=12592
        Reduce output records=1
        Spilled Records=25184
        Shuffled Maps =2
        Failed Shuffles=0
        Merged Map outputs=2
        GC time elapsed (ms)=380

```

```

CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=480247808
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=404717
File Output Format Counters
  Bytes Written=34
16/02/03 11:44:31 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_3_1_sorted_basket

```

3.3.1 Largest Basket: 37

Note: the largest basket of 37 occurred for many items, and not specific to the one item shown (it is simply the one that came out in the item name sorting)

```
In [4]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3_1_sorted_basket/part-*
| head -n 1
```

```

16/02/03 11:44:37 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
DAI35347      1060      0.00278343801861      37

```

3.3.1 Number of Unique Items: 12592

```
In [9]: !hdfs dfs -cat /user/konniam/week_03/hw_3_3_1/part-* | wc -l
```

```

16/02/01 11:15:13 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
12592

```

HW3.4. (Computationally prohibitive but then again Hadoop can handle this) Pairs

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a map-reduce program to find products which are frequently browsed together. Fix the support count (cooccurrence count) to $s = 100$ (i.e. product pairs need to occur together at least 100 times to be considered frequent) and find pairs of items (sometimes referred to itemsets of size 2 in association rule mining) that have a support count of 100 or more.

List the top 50 product pairs with corresponding support count (aka frequency), and relative frequency or support (number of records where they occur, the number of records where they occur/the number of baskets in the dataset) in decreasing order of support for frequent ($100 > \text{count}$) itemsets of size 2.

Use the Pairs pattern (lecture 3) to extract these frequent itemsets of size 2. Free free to use combiners if they bring value. Instrument your code with counters for count the number of times your mapper, combiner and reducers are called.

Please output records of the following form for the top 50 pairs (itemsets of size 2):

```
item1, item2, support count, support
```

Fix the ordering of the pairs lexicographically (left to right), and break ties in support (between pairs, if any exist) by taking the first ones in lexicographically increasing order.

Report the compute time for the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers) Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts.

Strategy

- 1st stage:
Mapper: basket \rightarrow (Item1, Item2, 1)
Reducer: (Item1, Item2, 1) \rightarrow (Item1, Item2, count, relative frequency)
(Use 2 mappers and 2 reducers)
- 2nd stage: Sort (Item1, Item2, count, relative frequency) by count

Mapper 3.4

```
In [60]: %%writefile mapper_3_4.py
#!/usr/bin/env python
# Mapper for 3.4
# Author: Konniam Chan
# Pairs analysis
import sys
from itertools import combinations

for line in sys.stdin:
    # Count how many mappers called
    sys.stderr.write("reporter:counter:MR-Counter,Num_mapper_calls,
1\n")
    record = line.strip().split()
    # Assume all items in basket co-occur, get all combinations of
items
    # Sort by alphanumerical order first and eliminate potential du
plicates
    record = sorted(set(record))
    for combo in combinations(record, 2):
        # Emit each co-occurrence pair
        print '%s\t%s' % (combo[0], combo[1], 1)
    # Assume 1 reducer in the next step
    # Emit total with key (*0 and *1) for order inversion (keep tra
ck of total number of baskets)
    # Assume output goes to 2 reducers
    print '%s\t%s' % ('*0' , 1)
    print '%s\t%s' % ('*1' , 1)
```

Overwriting mapper_3_4.py

Combiner 3.4

```
In [61]: %%writefile combiner_3_4.py
#!/usr/bin/env python
## Combiner for 3.4
# Author: Konniam Chan
# Pairs analysis
import sys
current_pair = None
current_count = 0

for line in sys.stdin:
    # Count how many combiners called
    sys.stderr.write("reporter:counter:MR-Counter,Num_combiner_calls,1\n")
    # Obtain pair and intermediate counts
    line = line.strip()
    pair, count = line.split('\t')
    count = int(count)
    # Rely on sorting to increment pair counts
    if current_pair == pair:
        current_count += count
    else:
        if current_pair:
            print '%s\t%s' % (current_pair, current_count)
        current_count = count
        current_pair = pair

# Output last pair
if current_pair == pair:
    print '%s\t%s' % (current_pair, current_count)
```

Overwriting combiner_3_4.py

Reducer 3.4

```
In [62]: %%writefile reducer_3_4.py
#!/usr/bin/env python
## Reducer for 3.4
# Author: Konniam Chan
# Pairs analysis
from __future__ import division
import sys
current_pair = None
current_count = 0
s = 100

for line in sys.stdin:
    # Count how many reducers called
    sys.stderr.write("reporter:counter:MR-Counter,Num_reducer_calls
,1\n")
    # Obtain pair and intermediate counts
    line = line.strip()
    pair, count = line.split('\t')
    count = int(count)
    # Rely on sorting to increment pair counts
    if current_pair == pair:
        current_count += count
    else:
        # Obtain total count and save to memory
        if current_pair in ['*0', '*1']:
            total = current_count
        elif current_pair:
            # Emit only if count is above s (100)
            if current_count >= s:
                item_1, item_2 = current_pair.split('.')
                print '%s\t%s\t%s\t%s' % (item_1, item_2, current_c
ount, current_count / total)
            current_count = count
            current_pair = pair

# Output last pair
if current_pair == pair:
    # Emit only if count is above s (100)
    if current_count >= s:
        item_1, item_2 = current_pair.split('.')
        print '%s\t%s\t%s\t%s' % (item_1, item_2, current_count, cu
rrent_count / total)
```

Overwriting reducer_3_4.py

```
In [63]: !chmod a+x *_3_4.py
```

3.4 MR 1st Stage

```
In [64]: # Shopping cart analysis with pairs, with counters for number of ma
ppers, combiners, and reducers
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_4
# Use 2 mappers and 2 reducers
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=2 \
-mapper $PWD/mapper_3_4.py \
-combiner $PWD/combiner_3_4.py \
-reducer $PWD/reducer_3_4.py \
-input /user/konniam/week_03/ProductPurchaseData.txt \
-output /user/konniam/week_03/hw_3_4
```

16/01/31 18:05:56 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

16/01/31 18:05:57 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.

Deleted /user/konniam/week_03/hw_3_4

16/01/31 18:05:59 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

packageJobJar: [/var/folders/l8/h5l_59852qscq403fs6q0xlh0000gn/T/ha
doo-unjar709333870281094859/] [] /var/folders/l8/h5l_59852qscq403f
s6q0xlh0000gn/T/streamjob7289565244085294957.jar tmpDir=null

16/01/31 18:06:00 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/01/31 18:06:01 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/01/31 18:06:02 INFO mapred.FileInputFormat: Total input paths to
process : 1

16/01/31 18:06:02 INFO mapreduce.JobSubmitter: number of splits:2

16/01/31 18:06:02 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454277284610_0016

16/01/31 18:06:02 INFO impl.YarnClientImpl: Submitted application a
pplication_1454277284610_0016

16/01/31 18:06:02 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454277284610
_0016/

16/01/31 18:06:02 INFO mapreduce.Job: Running job: job_145427728461
0_0016

16/01/31 18:06:11 INFO mapreduce.Job: Job job_1454277284610_0016 ru
nning in uber mode : false

16/01/31 18:06:11 INFO mapreduce.Job: map 0% reduce 0%

```

16/01/31 18:06:26 INFO mapreduce.Job: map 58% reduce 0%
16/01/31 18:06:28 INFO mapreduce.Job: map 67% reduce 0%
16/01/31 18:06:48 INFO mapreduce.Job: map 83% reduce 0%
16/01/31 18:06:49 INFO mapreduce.Job: map 100% reduce 0%
16/01/31 18:07:02 INFO mapreduce.Job: map 100% reduce 39%
16/01/31 18:07:04 INFO mapreduce.Job: map 100% reduce 81%
16/01/31 18:07:05 INFO mapreduce.Job: map 100% reduce 89%
16/01/31 18:07:06 INFO mapreduce.Job: map 100% reduce 92%
16/01/31 18:07:07 INFO mapreduce.Job: map 100% reduce 100%
16/01/31 18:07:08 INFO mapreduce.Job: Job job_1454277284610_0016 co
mpleted successfully

```

```

16/01/31 18:07:08 INFO mapreduce.Job: Counters: 52

```

File System Counters

```

FILE: Number of bytes read=22624207
FILE: Number of bytes written=45720068
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3462103
HDFS: Number of bytes written=51765
HDFS: Number of read operations=12
HDFS: Number of large read operations=0
HDFS: Number of write operations=4

```

Job Counters

```

Launched map tasks=2
Launched reduce tasks=2
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)
=70318
Total time spent by all reduces in occupied slots (
ms)=32606
Total time spent by all map tasks (ms)=70318
Total time spent by all reduce tasks (ms)=32606
Total vcore-seconds taken by all map tasks=70318
Total vcore-seconds taken by all reduce tasks=32606
Total megabyte-seconds taken by all map tasks=72005
632
Total megabyte-seconds taken by all reduce tasks=33
388544

```

Map-Reduce Framework

```

Map input records=31101
Map output records=2596216
Map output bytes=50991290
Map output materialized bytes=22624219
Input split bytes=236
Combine input records=2596216
Combine output records=1026711
Reduce input groups=877097
Reduce shuffle bytes=22624219
Reduce input records=1026711

```



```

Reduce output records=1334
Spilled Records=2053422
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=932
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=599785472
MR-Counter
  Num_combiner_calls=2596216
  Num_mapper_calls=31101
  Num_reducer_calls=1026711
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=3461867
File Output Format Counters
  Bytes Written=51765
16/01/31 18:07:08 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_4

```

```
In [65]: !hdfs dfs -cat /user/konniam/week_03/hw_3_4/part* | wc -l
```

```

16/01/31 18:07:15 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
1334

```

3.4 MR 2nd Stage (sort by count)

```
In [5]: # Shopping cart analysis with pairs, with counters for number of ma
ppers, combiners, and reducers
# Sort output of 1st stage MR to get top support counts
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_4_sorted
# Use 2 mappers and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educe.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=3 \
-D mapreduce.partition.keycomparator.options='-k3,3nr -k1,2' \
-mapper /bin/cat \
-reducer $PWD/reducer_top50.py \
-input /user/konniam/week_03/hw_3_4/part* \
-output /user/konniam/week_03/hw_3_4_sorted
```

```
16/02/03 11:45:23 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 11:45:24 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_4_sorted
16/02/03 11:45:25 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/ha
doo-unjar3029703903248871676/] [] /var/folders/l8/h51_59852qscq403
fs6q0xlh0000gn/T/streamjob984625331549406897.jar tmpDir=null
16/02/03 11:45:27 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:45:27 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 11:45:28 INFO mapred.FileInputFormat: Total input paths to
process : 2
16/02/03 11:45:29 INFO mapreduce.JobSubmitter: number of splits:2
16/02/03 11:45:29 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0026
16/02/03 11:45:29 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0026
16/02/03 11:45:29 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0026/
16/02/03 11:45:29 INFO mapreduce.Job: Running job: job_145448753435
8_0026
16/02/03 11:45:38 INFO mapreduce.Job: Job job_1454487534358_0026 ru
nning in uber mode : false
```

```

16/02/03 11:45:38 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:45:50 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 11:45:57 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 11:45:58 INFO mapreduce.Job: Job job_1454487534358_0026 co
mpleted successfully
16/02/03 11:45:58 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=54439
        FILE: Number of bytes written=462868
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=51989
        HDFS: Number of bytes written=1898
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=17244
        Total time spent by all reduces in occupied slots (
ms)=5255
        Total time spent by all map tasks (ms)=17244
        Total time spent by all reduce tasks (ms)=5255
        Total vcore-seconds taken by all map tasks=17244
        Total vcore-seconds taken by all reduce tasks=5255
        Total megabyte-seconds taken by all map tasks=17657
856
        Total megabyte-seconds taken by all reduce tasks=53
81120
    Map-Reduce Framework
        Map input records=1334
        Map output records=1334
        Map output bytes=51765
        Map output materialized bytes=54445
        Input split bytes=224
        Combine input records=0
        Combine output records=0
        Reduce input groups=1334
        Reduce shuffle bytes=54445
        Reduce input records=1334
        Reduce output records=50
        Spilled Records=2668
        Shuffled Maps =2
        Failed Shuffles=0
        Merged Map outputs=2
        GC time elapsed (ms)=360

```

```

CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=493355008
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=51765
File Output Format Counters
  Bytes Written=1898
16/02/03 11:45:58 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_4_sorted

```

3.4 Results

Top 50 product pairs by support counts

```
In [6]: !hdfs dfs -cat /user/konniam/week_03/hw_3_4_sorted/part*
```

```

16/02/03 11:46:03 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

```

DAI62779	ELE17451	1592	0.0511880646925
FRO40251	SNA80324	1412	0.0454004694383
DAI75645	FRO40251	1254	0.0403202469374
FRO40251	GRO85051	1213	0.0390019613517
DAI62779	GRO73461	1139	0.0366226166361
DAI75645	SNA80324	1130	0.0363332368734
DAI62779	FRO40251	1070	0.0344040384554
DAI62779	SNA80324	923	0.0296775023311
DAI62779	DAI85309	918	0.0295167357963
ELE32164	GRO59710	911	0.0292916626475
DAI62779	DAI75645	882	0.0283592167454
FRO40251	GRO73461	882	0.0283592167454
DAI62779	ELE92920	877	0.0281984502106
FRO40251	FRO92469	835	0.026848011318
DAI62779	ELE32164	832	0.0267515513971
DAI75645	GRO73461	712	0.0228931545609
DAI43223	ELE32164	711	0.022861001254
DAI62779	GRO30386	709	0.02279669464
ELE17451	FRO40251	697	0.0224108549564
DAI85309	ELE99737	659	0.0211890292917
DAI62779	ELE26917	650	0.020899649529

GRO21487	GRO73461	631	0.0202887366966
DAI62779	SNA45677	604	0.0194205974084
ELE17451	SNA80324	597	0.0191955242597
DAI62779	GRO71621	595	0.0191312176457
DAI62779	SNA55762	593	0.0190669110318
DAI62779	DAI83733	586	0.018841837883
ELE17451	GRO73461	580	0.0186489180412
GRO73461	SNA80324	562	0.0180701585158
DAI62779	GRO59710	561	0.0180380052088
DAI62779	FRO80039	550	0.0176843188322
DAI75645	ELE17451	547	0.0175878589113
DAI62779	SNA93860	537	0.0172663258416
DAI55148	DAI62779	526	0.016912639465
DAI43223	GRO59710	512	0.0164624931674
ELE17451	ELE32164	511	0.0164303398605
DAI62779	SNA18336	506	0.0162695733256
ELE32164	GRO73461	486	0.0156265071863
DAI62779	FRO78087	482	0.0154978939584
DAI85309	ELE17451	482	0.0154978939584
DAI62779	GRO94758	479	0.0154014340375
DAI62779	GRO21487	471	0.0151442075817
GRO85051	SNA80324	471	0.0151442075817
ELE17451	GRO30386	468	0.0150477476608
FRO85978	SNA95666	463	0.014886981126
DAI62779	FRO19221	462	0.014854827819
DAI62779	GRO46854	461	0.0148226745121
DAI43223	DAI62779	459	0.0147583678981
ELE92920	SNA18336	455	0.0146297546703
DAI88079	FRO40251	446	0.0143403749076

3.4 Pairs Statistics

Property	Configuration
Operating System	Mac
Cores	2
Memory	4GB
Number of Tuples Processed	
- Mappers	31,101
- Combiners	2,596,216
- Reducers	1,026,711

Stage	Elapsed Time (s)
1st Stage MapReduce (2 mappers, 2 reducers, with combiners)	56
2nd Stage MapReduce (2 mappers, 1 reducer)	19
Total	75

HW3.5: Stripes

Repeat 3.4 using the stripes design pattern for finding cooccurring pairs.

Report the compute times for stripes job versus the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts. Discuss the differences in these counts between the Pairs and Stripes jobs

Strategy

- 1st stage:
Mapper: basket -> (Item_i, Stripe_i) [where Stripe_i is a json formatted string]
Reducer: (Item_i, Stripe_i) -> (Item1, Item2, count, relative frequency)
(Use 2 mappers and 2 reducers)
- 2nd stage: Sort (Item1, Item2, count, relative frequency) by count

3.5 Mapper

```
In [8]: %%writefile mapper_3_5.py
#!/usr/bin/env python
# Mapper for 3.5
# Author: Konniam Chan
# Stripes shopping cart analysis
import sys
import json
from itertools import combinations

for line in sys.stdin:
    # Count how many mappers called
    sys.stderr.write("reporter:counter:MR-Counter,Num_mapper_calls,
1\n")
    # Read record
    record = line.strip().split()
    # Assume all items in basket co-occur, get all combinations of
items
    # Sort by alphanumerical order first and eliminate potential du
plicates
    record = sorted(set(record))
    # Skip last item because it will always be part of another item
's row
    stripes = {item: {} for item in record[:-1]}
    # Create stripes
    for pair in combinations(record, 2):
        stripes[pair[0]][pair[1]] = 1
    # Emit each stripe
    for k,v in stripes.items():
        print '%s\t%s' % (k, json.dumps(v))
    # Emit basket count for order inversion
    # Assume 2 reducers
    print '%s\t%s' % ('*0', json.dumps({"*0": 1}))
    print '%s\t%s' % ('*1', json.dumps({"*1": 1}))
```

Overwriting mapper_3_5.py

3.5 Combiner

```
In [7]: %%writefile combiner_3_5.py
#!/usr/bin/env python
# Combiner for 3.5
# Author: Konniam Chan
# Stripes shopping cart analysis
import sys
import json

current_item = None
current_stripe = {}

# Function to add stripes
def add_stripes(current_stripe, stripe):
    for key in stripe:
        if key in current_stripe:
            current_stripe[key] += stripe[key]
        else:
            current_stripe[key] = stripe[key]
    return current_stripe

for line in sys.stdin:
    # Count how many combiners called
    sys.stderr.write("reporter:counter:MR-Counter,Num_combiner_calls,1\n")
    # Obtain stripe
    line = line.strip()
    item, stripe = line.split('\t')
    # Read json string, then convert to Counter dictionary
    stripe = json.loads(stripe)
    # Rely on sorting to collect stripes
    if current_item == item:
        current_stripe = add_stripes(current_stripe, stripe)
    else:
        if current_item:
            print '%s\t%s' % (current_item, json.dumps(current_stripe))
            current_stripe = stripe
            current_item = item

# Output last stripe
if current_item == item:
    print '%s\t%s' % (current_item, json.dumps(current_stripe))
```

Overwriting combiner_3_5.py

3.5 Reducer

```
In [19]: %%writefile reducer_3_5.py
```



```
#!/usr/bin/env python
## Reducer for 3.5
# Author: Konniam Chan
# Stripes shopping cart analysis
from __future__ import division
import sys
import json

current_item = None
current_stripe = {}
s = 100

# Function to add stripes
def add_stripes(current_stripe, stripe):
    for key in stripe:
        if key in current_stripe:
            current_stripe[key] += stripe[key]
        else:
            current_stripe[key] = stripe[key]
    return current_stripe

for line in sys.stdin:
    # Count how many reducers called
    sys.stderr.write("reporter:counter:MR-Counter,Num_reducer_calls
,1\n")
    # Obtain stripe
    line = line.strip()
    item, stripe = line.split('\t')
    # Read json string, then convert to Counter dictionary
    stripe = json.loads(stripe)
    # Rely on sorting to collect stripes
    if current_item == item:
        current_stripe = add_stripes(current_stripe, stripe)
    else:
        # Obtain total count and save to memory
        if current_item == '*0':
            total = current_stripe['*0']
        elif current_item == '*1':
            total = current_stripe['*1']
        elif current_item:
            # Output all neighbor pairs in the stripe
            for neighbor, current_count in current_stripe.items():
                # Output in support is at least 100
                if current_count >= s:
                    print '%s\t%s\t%s\t%s' % (current_item, neighbor,
current_count, current_count/total)
            current_stripe = stripe
            current_item = item

# Output last stripe
```

```

if current_item == item:
    for neighbor, current_count in current_stripe.items():
        if current_count >= s:
            print '%s\t%s\t%s\t%s' % (current_item, neighbor, current_count, current_count/total)

```

Overwriting reducer_3_5.py

In [79]: `!chmod a+x *_3_5.py`

3.5 MR 1st Stage

In [20]: `# Shopping cart analysis with stripes, with counters for number of mappers, combiners, and reducers`
`!hdfs dfs -rm -r /user/konniam/week_03/hw_3_5`
`# Use 2 mappers and 2 reducers`
`!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.jar \`
`-D mapreduce.job.maps=2 \`
`-D mapreduce.job.reduces=2 \`
`-mapper $PWD/mapper_3_5.py \`
`-reducer $PWD/reducer_3_5.py \`
`-combiner $PWD/combiner_3_5.py \`
`-input /user/konniam/week_03/ProductPurchaseData.txt \`
`-output /user/konniam/week_03/hw_3_5`

```

16/02/03 11:58:50 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
16/02/03 11:58:51 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /user/konniam/week_03/hw_3_5
16/02/03 11:58:52 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/hadoop-unjar5986202487866547422/] [] /var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/streamjob1043533591408109191.jar tmpDir=null
16/02/03 11:58:54 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
16/02/03 11:58:54 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
16/02/03 11:58:55 INFO mapred.FileInputFormat: Total input paths to process : 1
16/02/03 11:58:55 INFO mapreduce.JobSubmitter: number of splits:2
16/02/03 11:58:56 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1454487534358_0027
16/02/03 11:58:56 INFO impl.YarnClientImpl: Submitted application a

```

```

pplication_1454487534358_0027
16/02/03 11:58:56 INFO mapreduce.Job: The url to track the job: http://Konniams-MacBook-Air.local:8088/proxy/application_1454487534358_0027/
16/02/03 11:58:56 INFO mapreduce.Job: Running job: job_1454487534358_0027
16/02/03 11:59:04 INFO mapreduce.Job: Job job_1454487534358_0027 running in uber mode : false
16/02/03 11:59:04 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 11:59:21 INFO mapreduce.Job:  map 53% reduce 0%
16/02/03 11:59:24 INFO mapreduce.Job:  map 67% reduce 0%
16/02/03 11:59:34 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 11:59:45 INFO mapreduce.Job:  map 100% reduce 50%
16/02/03 11:59:46 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 11:59:47 INFO mapreduce.Job: Job job_1454487534358_0027 completed successfully
16/02/03 11:59:47 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=15677786
        FILE: Number of bytes written=31827226
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3462103
        HDFS: Number of bytes written=51765
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=54133
        Total time spent by all reduces in occupied slots (ms)=18186
        Total time spent by all map tasks (ms)=54133
        Total time spent by all reduce tasks (ms)=18186
        Total vcore-seconds taken by all map tasks=54133
        Total vcore-seconds taken by all reduce tasks=18186
        Total megabyte-seconds taken by all map tasks=55432192
        Total megabyte-seconds taken by all reduce tasks=18622464
    Map-Reduce Framework
        Map input records=31101
        Map output records=411922
        Map output bytes=42450230
        Map output materialized bytes=15677798
        Input split bytes=236

```

```

Combine input records=411922
Combine output records=16944
Reduce input groups=12013
Reduce shuffle bytes=15677798
Reduce input records=16944
Reduce output records=1334
Spilled Records=33888
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=662
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=586153984
MR-Counter
  Num_combiner_calls=411922
  Num_mapper_calls=31101
  Num_reducer_calls=16944
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=3461867
File Output Format Counters
  Bytes Written=51765
16/02/03 11:59:47 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_5

```

```
In [21]: !hdfs dfs -cat /user/konniam/week_03/hw_3_5/part* | wc -l
```

```

16/02/03 12:05:17 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
1334

```

3.5 MR 2nd Stage (sort by count)

```
In [22]: # Shopping cart analysis with pairs, with counters for number of ma
ppers, combiners, and reducers
# Sort output of 1st stage MR to get top support counts
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_5_sorted
# Use 2 mappers and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educe.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=3 \
-D mapreduce.partition.keycomparator.options='-k3,3nr -k1,2' \
-mapper /bin/cat \
-reducer $PWD/reducer_top50.py \
-input /user/konniam/week_03/hw_3_5/part* \
-output /user/konniam/week_03/hw_3_5_sorted
```

```
16/02/03 12:05:42 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 12:05:43 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_5_sorted
16/02/03 12:05:45 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar784545658839887956/] [] /var/folders/l8/h51_59852qscq403f
s6q0xlh0000gn/T/streamjob7683328440605710148.jar tmpDir=null
16/02/03 12:05:47 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 12:05:47 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 12:05:48 INFO mapred.FileInputFormat: Total input paths to
process : 2
16/02/03 12:05:48 INFO mapreduce.JobSubmitter: number of splits:3
16/02/03 12:05:49 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0028
16/02/03 12:05:49 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0028
16/02/03 12:05:50 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0028/
16/02/03 12:05:50 INFO mapreduce.Job: Running job: job_145448753435
8_0028
16/02/03 12:05:59 INFO mapreduce.Job: Job job_1454487534358_0028 ru
nning in uber mode : false
```

```

16/02/03 12:05:59 INFO mapreduce.Job: map 0% reduce 0%
16/02/03 12:06:13 INFO mapreduce.Job: map 100% reduce 0%
16/02/03 12:06:21 INFO mapreduce.Job: map 100% reduce 100%
16/02/03 12:06:22 INFO mapreduce.Job: Job job_1454487534358_0028 completed successfully

```

```

16/02/03 12:06:23 INFO mapreduce.Job: Counters: 49

```

File System Counters

```

FILE: Number of bytes read=54439
FILE: Number of bytes written=580891
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=54891
HDFS: Number of bytes written=1898
HDFS: Number of read operations=12
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

```

Job Counters

```

Launched map tasks=3
Launched reduce tasks=1
Data-local map tasks=3
Total time spent by all maps in occupied slots (ms)
=35198
Total time spent by all reduces in occupied slots (
ms)=5756
Total time spent by all map tasks (ms)=35198
Total time spent by all reduce tasks (ms)=5756
Total vcore-seconds taken by all map tasks=35198
Total vcore-seconds taken by all reduce tasks=5756
Total megabyte-seconds taken by all map tasks=36042
752
Total megabyte-seconds taken by all reduce tasks=58
94144

```

Map-Reduce Framework

```

Map input records=1334
Map output records=1334
Map output bytes=51765
Map output materialized bytes=54451
Input split bytes=336
Combine input records=0
Combine output records=0
Reduce input groups=1334
Reduce shuffle bytes=54451
Reduce input records=1334
Reduce output records=50
Spilled Records=2668
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=517

```

```

CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=685244416
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=54555
File Output Format Counters
  Bytes Written=1898
16/02/03 12:06:23 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_5_sorted

```

3.5 Results

Top 50 product pairs by support counts

```
In [23]: !hdfs dfs -cat /user/konniam/week_03/hw_3_5_sorted/part*
```

```

16/02/03 12:06:30 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

```

DAI62779	ELE17451	1592	0.0511880646925
FRO40251	SNA80324	1412	0.0454004694383
DAI75645	FRO40251	1254	0.0403202469374
FRO40251	GRO85051	1213	0.0390019613517
DAI62779	GRO73461	1139	0.0366226166361
DAI75645	SNA80324	1130	0.0363332368734
DAI62779	FRO40251	1070	0.0344040384554
DAI62779	SNA80324	923	0.0296775023311
DAI62779	DAI85309	918	0.0295167357963
ELE32164	GRO59710	911	0.0292916626475
DAI62779	DAI75645	882	0.0283592167454
FRO40251	GRO73461	882	0.0283592167454
DAI62779	ELE92920	877	0.0281984502106
FRO40251	FRO92469	835	0.026848011318
DAI62779	ELE32164	832	0.0267515513971
DAI75645	GRO73461	712	0.0228931545609
DAI43223	ELE32164	711	0.022861001254
DAI62779	GRO30386	709	0.02279669464
ELE17451	FRO40251	697	0.0224108549564
DAI85309	ELE99737	659	0.0211890292917
DAI62779	ELE26917	650	0.020899649529

GRO21487	GRO73461	631	0.0202887366966
DAI62779	SNA45677	604	0.0194205974084
ELE17451	SNA80324	597	0.0191955242597
DAI62779	GRO71621	595	0.0191312176457
DAI62779	SNA55762	593	0.0190669110318
DAI62779	DAI83733	586	0.018841837883
ELE17451	GRO73461	580	0.0186489180412
GRO73461	SNA80324	562	0.0180701585158
DAI62779	GRO59710	561	0.0180380052088
DAI62779	FRO80039	550	0.0176843188322
DAI75645	ELE17451	547	0.0175878589113
DAI62779	SNA93860	537	0.0172663258416
DAI55148	DAI62779	526	0.016912639465
DAI43223	GRO59710	512	0.0164624931674
ELE17451	ELE32164	511	0.0164303398605
DAI62779	SNA18336	506	0.0162695733256
ELE32164	GRO73461	486	0.0156265071863
DAI62779	FRO78087	482	0.0154978939584
DAI85309	ELE17451	482	0.0154978939584
DAI62779	GRO94758	479	0.0154014340375
DAI62779	GRO21487	471	0.0151442075817
GRO85051	SNA80324	471	0.0151442075817
ELE17451	GRO30386	468	0.0150477476608
FRO85978	SNA95666	463	0.014886981126
DAI62779	FRO19221	462	0.014854827819
DAI62779	GRO46854	461	0.0148226745121
DAI43223	DAI62779	459	0.0147583678981
ELE92920	SNA18336	455	0.0146297546703
DAI88079	FRO40251	446	0.0143403749076

3.5 Stripes Statistics

Property	Configuration
Operating System	Mac
Cores	2
Memory	4GB
Number of Tuples Processed	
- Mappers	31,101
- Combiners	411,922
- Reducers	16,944

Stage	Elapsed Time (s)
1st Stage MapReduce (2 mappers, 2 reducers, with combiners)	41
2nd Stage MapReduce (2 mappers, 1 reducer)	22
Total	63

3.5 Response

In the pairs formulation, there are many records transmitted between the mappers and the reducers. In this particular problem, the reducer received over 1,000,000 records. In the stripes formulation, due to the more condensed structure of the stripe, there are only about 17,000 records. In jobs where there are a large number of mappers and reducers on different networks, the number of tuples sent in the pairs approach could be prohibitive. In our pseudo-distributed environment, both mappers and reducers live on the same network and we didn't incur any network costs. The stripes approach is faster by 12 seconds. If the network costs were more significant, we would see stripes performing significantly better.

HW3.6

What is the Apriori algorithm? Describe an example use in your domain of expertise and what kind of . Define confidence and lift.

Apriori

The Apriori algorithm is used to find frequent item sets. These sets are useful in generating association rules, which are widely used in recommendations. The central idea is that if a set of items A is frequent, so is any subset of A. Thus, if any subset of A is not frequent, then A cannot be frequent. This idea is used to limit the large space of higher-number item sets.

The algorithm works in stages. In the first pass, we look at all baskets of transactions and generates a list of counts for each item (C1). We filter it based on a specified minimum support count and obtain a set of frequent items of size 1 (L1). In the second pass, we look at all baskets, filtered out all items NOT in L1, then use the rest to generate candidate paris (C2). We filter this list to generate frequent items of size 2 (L2). In the third pass, we look at all baskets and filtered out all items not in L1. We form pairs, then filter out items not in L2. We then form triplet candidates (C3), filter to get L3. Repeat until we get to the size of frequent items that we want.

Confidence and Lift

For an association rule $A \rightarrow b$, **confidence** measures how "likely" b would occur, in the baskets that contain A. High confidence means that if A is in the basket, b is likley to be in the same basket. The formula is: $\frac{\text{support}(A \cup b)}{\text{support}(A)}$.

Lift measures the ratio between confidence and how likley is b to occur unconditionally. It can be calculated as $\frac{\text{Confidence}(A \rightarrow b)}{\text{Confidence}(\text{null} \rightarrow b)}$ where $\text{Confidence}(\text{null} \rightarrow b) = \text{support}(b)$

Applications

One could use Apriori algorithm to create marketing campaigns in retail stores, let's say at Starbucks. If an association rule shows that two food items are often purchased together, we can advertise both items in some fashion to increase revenue per transaction.

HW3.7. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a program using the A-priori algorithm to find products which are frequently browsed together. Fix the support to $s = 100$ (i.e. product sets need to occur together at least 100 times to be considered frequent) and find itemsets of size 2 and 3.

Then extract association rules from these frequent items.

A rule is of the form:

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}$.

List the top 10 discovered rules in decreasing order of confidence in the following format

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}$, supportCount ,support, confidence

Strategy

We will do 3 passes over the data, to get the set of triples with support of at least 100 (L3). Then we can easily do another MR job to calculate confidence by dividing L3 counts with the corresponding L2 counts. Thus, in total we will run 4 MR jobs.

1st pass of Apriori Algorithm

```
In [2]: %%writefile mapper_3_7_pass1.py
#!/usr/bin/env python
# Mapper for 3.7
# Author: Konniam Chan
# Apriori Algorithm
import sys

for line in sys.stdin:
    record = line.strip().split()
    # Emit count of item
    for item in record:
        print '%s\t%s' % (item, 1)
```

Writing mapper_3_7_pass1.py

```
In [25]: %%writefile combiner_3_7_pass1.py
#!/usr/bin/env python
## Combiner for 3.7
# Author: Konniam Chan
# Apriori
import sys
current_item = None
current_count = 0
item = None

for line in sys.stdin:
    # Obtain item and intermediate counts
    line = line.strip()
    item, count = line.split('\t')
    count = int(count)
    # Rely on sorting to increment item counts
    if current_item == item:
        current_count += count
    else:
        if current_item:
            print '%s\t%s' % (current_item, current_count)
            current_count = count
            current_item = item

# Output last item
if current_item == item:
    print '%s\t%s' % (current_item, current_count)
```

Writing combiner_3_7_pass1.py

```
In [4]: %%writefile reducer_3_7_pass1.py
#!/usr/bin/env python
## Reducer for 3.7
# Author: Konniam Chan
# Apriori
import sys
current_item = None
current_count = 0
item = None
SUPPORT = 100

for line in sys.stdin:
    # Obtain item and intermediate counts
    line = line.strip()
    item, count = line.split('\t')
    count = int(count)
    # Rely on sorting to increment item counts
    if current_item == item:
        current_count += count
    else:
        if current_item and current_count >= SUPPORT:
            print '%s\t%s' % (current_item, current_count)
        current_count = count
        current_item = item

# Output last item
if current_item == item and current_count >= SUPPORT:
    print '%s\t%s' % (current_item, current_count)
```

Writing reducer_3_7_pass1.py

```
In [5]: !chmod a+x *_3_7_pass1.py
```

Obtain L1

```
In [6]: # Apriori
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_7_pass1
# Use 2 mappers and 2 reducers
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-mapper $PWD/mapper_3_7_pass1.py \
-combiner $PWD/combiner_3_7_pass1.py \
-reducer $PWD/reducer_3_7_pass1.py \
-input /user/konniam/week_03/ProductPurchaseData.txt \
-output /user/konniam/week_03/hw_3_7_pass1
```

```
16/02/03 23:28:15 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
16/02/03 23:28:16 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.
Deleted /user/konniam/week_03/hw_3_7_pass1
16/02/03 23:28:18 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
packageJobJar: [/var/folders/l8/h5l_59852qscq403fs6q0xlh0000gn/T/ha
doop-unjar4403250698977015195/] [] /var/folders/l8/h5l_59852qscq403
fs6q0xlh0000gn/T/streamjob7955926000752034618.jar tmpDir=null
16/02/03 23:28:20 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 23:28:21 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/03 23:28:23 INFO mapred.FileInputFormat: Total input paths to
process : 1
16/02/03 23:28:23 INFO mapreduce.JobSubmitter: number of splits:2
16/02/03 23:28:23 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0031
16/02/03 23:28:24 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0031
16/02/03 23:28:24 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0031/
16/02/03 23:28:24 INFO mapreduce.Job: Running job: job_145448753435
8_0031
16/02/03 23:28:38 INFO mapreduce.Job: Job job_1454487534358_0031 ru
nning in uber mode : false
16/02/03 23:28:38 INFO mapreduce.Job:  map 0% reduce 0%
16/02/03 23:28:54 INFO mapreduce.Job:  map 100% reduce 0%
16/02/03 23:29:04 INFO mapreduce.Job:  map 100% reduce 100%
16/02/03 23:29:06 INFO mapreduce.Job: Job job_1454487534358_0031 co
mpleted successfully
16/02/03 23:29:06 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=235933
        FILE: Number of bytes written=825667
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3462103
        HDFS: Number of bytes written=8471
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
```

```

Launched reduce tasks=1
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)
=27969
Total time spent by all reduces in occupied slots (
ms)=6184
Total time spent by all map tasks (ms)=27969
Total time spent by all reduce tasks (ms)=6184
Total vcore-seconds taken by all map tasks=27969
Total vcore-seconds taken by all reduce tasks=6184
Total megabyte-seconds taken by all map tasks=28640
256
Total megabyte-seconds taken by all reduce tasks=63
32416

```

Map-Reduce Framework

```

Map input records=31101
Map output records=380824
Map output bytes=4189064
Map output materialized bytes=235939
Input split bytes=236
Combine input records=380824
Combine output records=17745
Reduce input groups=12592
Reduce shuffle bytes=235939
Reduce input records=17745
Reduce output records=647
Spilled Records=35490
Shuffled Maps =2
Failed Shuffles=0
Merged Map outputs=2
GC time elapsed (ms)=423
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=487063552

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=3461867
```

File Output Format Counters

```
Bytes Written=8471
```

```

16/02/03 23:29:06 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_7_pass1

```

```
In [7]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass1/part* | wc -l
```

```
16/02/03 23:30:17 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
        647
```

```
In [8]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass1/part* | sort -k2,
2nr | head
```

```
16/02/03 23:30:20 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
DAI62779          6667
FRO40251          3881
ELE17451          3875
GRO73461          3602
SNA80324          3044
ELE32164          2851
DAI75645          2736
SNA45677          2455
FRO31317          2330
DAI85309          2293
```

```
In [9]: !hdfs dfs -getmerge /user/konniam/week_03/hw_3_7_pass1 L1_counts.tx
t
```

```
16/02/03 23:30:24 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

2nd Pass to Obtain L2


```
In [10]: %%writefile mapper_3_7_pass2.py
#!/usr/bin/env python
# Mapper for 3.7
# Author: Konniam Chan
# Apriori Algorithm
import sys
from itertools import combinations

# Array to hold frequent items
freq_items_L1 = []
with open("L1_counts.txt", "r") as f:
    for line in f:
        item, _ = line.strip().split()
        freq_items_L1.append(item)

# Accept input from Hadoop streaming
for line in sys.stdin:
    items = line.strip().split()
    # Filter to only include items in frequent list
    items = filter(lambda x: x in freq_items_L1, items)
    # Sort items
    items = sorted(set(items))
    # Create candidate pairs (C2)
    for pair in combinations(items, 2):
        print '%s.%s\t%s' % (pair[0], pair[1], 1)
```

Writing mapper_3_7_pass2.py

```
In [11]: !cp combiner_3_7_pass1.py combiner_3_7_pass2.py
!cp reducer_3_7_pass1.py reducer_3_7_pass2.py
```

```
In [12]: !chmod a+x *_3_7_pass2.py
```

Obtain L2

```
In [13]: # Apriori
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_7_pass2
# Use 2 mappers and 2 reducers
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=2 \
-mapper $PWD/mapper_3_7_pass2.py \
-combiner $PWD/combiner_3_7_pass2.py \
-reducer $PWD/reducer_3_7_pass2.py \
-input /user/konniam/week_03/ProductPurchaseData.txt \
-output /user/konniam/week_03/hw_3_7_pass2 \
-file L1_counts.txt
```

16/02/03 23:30:58 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

16/02/03 23:30:59 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.

Deleted /user/konniam/week_03/hw_3_7_pass2

16/02/03 23:31:01 WARN streaming.StreamJob: -file option is depreca
ted, please use generic option -files instead.

16/02/03 23:31:01 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

packageJobJar: [L1_counts.txt, /var/folders/l8/h5l_59852qscq403fs6q
0xlh0000gn/T/hadoop-unjar721548109007871352/] [] /var/folders/l8/h5
l_59852qscq403fs6q0xlh0000gn/T/streamjob6445342460718776132.jar tmp
Dir=null

16/02/03 23:31:03 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/02/03 23:31:03 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/02/03 23:31:05 INFO mapred.FileInputFormat: Total input paths to
process : 1

16/02/03 23:31:05 INFO mapreduce.JobSubmitter: number of splits:2

16/02/03 23:31:05 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0032

16/02/03 23:31:06 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0032

16/02/03 23:31:06 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0032/

16/02/03 23:31:06 INFO mapreduce.Job: Running job: job_145448753435
8_0032

16/02/03 23:31:17 INFO mapreduce.Job: Job job_1454487534358_0032 ru
nning in uber mode : false

16/02/03 23:31:17 INFO mapreduce.Job: map 0% reduce 0%

```

16/02/03 23:31:34 INFO mapreduce.Job: map 51% reduce 0%
16/02/03 23:31:37 INFO mapreduce.Job: map 67% reduce 0%
16/02/03 23:31:45 INFO mapreduce.Job: map 83% reduce 0%
16/02/03 23:31:46 INFO mapreduce.Job: map 100% reduce 0%
16/02/03 23:32:00 INFO mapreduce.Job: map 100% reduce 50%
16/02/03 23:32:02 INFO mapreduce.Job: map 100% reduce 100%
16/02/03 23:32:02 INFO mapreduce.Job: Job job_1454487534358_0032 co
mpleted successfully
16/02/03 23:32:02 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=5030492
        FILE: Number of bytes written=10537330
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3462103
        HDFS: Number of bytes written=29355
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)
=49972
        Total time spent by all reduces in occupied slots (
ms)=26931
        Total time spent by all map tasks (ms)=49972
        Total time spent by all reduce tasks (ms)=26931
        Total vcore-seconds taken by all map tasks=49972
        Total vcore-seconds taken by all reduce tasks=26931
        Total megabyte-seconds taken by all map tasks=51171
328
        Total megabyte-seconds taken by all reduce tasks=27
577344
    Map-Reduce Framework
        Map input records=31101
        Map output records=1358592
        Map output bytes=27171840
        Map output materialized bytes=5030504
        Input split bytes=236
        Combine input records=1358592
        Combine output records=227163
        Reduce input groups=149097
        Reduce shuffle bytes=5030504
        Reduce input records=227163
        Reduce output records=1334
        Spilled Records=454326
        Shuffled Maps =4

```

```

Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=569
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=581959680

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=3461867
File Output Format Counters
  Bytes Written=29355
16/02/03 23:32:02 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_7_pass2

```

```
In [14]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass2/part* | wc -l
```

```

16/02/03 23:32:08 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
1334

```

```
In [15]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass2/part* | sort -k2,
2nr | head
```

```

16/02/03 23:32:12 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
DAI62779.ELE17451      1592
FRO40251.SNA80324      1412
DAI75645.FRO40251      1254
FRO40251.GRO85051      1213
DAI62779.GRO73461      1139
DAI75645.SNA80324      1130
DAI62779.FRO40251      1070
DAI62779.SNA80324       923
DAI62779.DAI85309       918
ELE32164.GRO59710       911

```

```
In [16]: !hdfs dfs -getmerge /user/konniam/week_03/hw_3_7_pass2 L2_counts.txt
```

```
16/02/03 23:32:16 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
```

Third pass to obtain L3

```
In [29]: %%writefile mapper_3_7_pass3.py
#!/usr/bin/env python
# Mapper for 3.7
# Author: Konniam Chan
# Apriori Algorithm
import sys
from itertools import combinations

# Arrays to hold frequent items
freq_items_L1 = []
with open("L1_counts.txt", "r") as f:
    for line in f:
        item, _ = line.strip().split()
        freq_items_L1.append(item)

freq_items_L2 = []
with open("L2_counts.txt", "r") as f:
    for line in f:
        item, _ = line.strip().split()
        freq_items_L2.append(item)

# Function to yield a pair for evaluation
def generate_pairs(items):
    '''
    Input: list of items filtered by L1
    Output: pairs of items, filtered by L2
    '''
    for pair in combinations(items, 2):
        if '.'.join([pair[0], pair[1]]) in freq_items_L2:
            yield pair

# Accept input from Hadoop streaming
for line in sys.stdin:
    items = line.strip().split()
    # Filter using L1 and sort
    items = filter(lambda x: x in freq_items_L1, items)
    items = sorted(set(items))
    # Form C3 from filtered L2 pairs
    for pair in generate_pairs(items):
        first, second = pair[0], pair[1]
```

```
        second_index = items.index(second)
        # Only look at items after the pair to find match (avoid double counting)
        for third in items[second_index+1:]:
            # Check if triple is a viable candidate
            if '.'.join([second, third]) in freq_items_L2 and '.'.join([first, third]) in freq_items_L2:
                # Triple is possible
                print '%s.%s.%s\t%s' % (first, second, third, 1)
        # Emit total for order inversion
        print '%s\t%s' % ('*0', 1)
        print '%s\t%s' % ('*1', 1)
```

Overwriting mapper_3_7_pass3.py

```
In [26]: # Use same combiner as pass 1
!mv combiner_3_7_pass1.py combiner_3_7_pass3.py
```

```
In [37]: %%writefile reducer_3_7_pass3.py
#!/usr/bin/env python
## Reducer for 3.7
# Author: Konniam Chan
# Apriori
from __future__ import division
import sys
current_triple = None
current_count = 0
triple = None
SUPPORT = 100

for line in sys.stdin:
    # Obtain triple and intermediate counts
    line = line.strip()
    triple, count = line.split('\t')
    count = int(count)
    # Rely on sorting to increment triple counts
    if current_triple == triple:
        current_count += count
    else:
        if current_triple in ['*0', '*1']:
            total = current_count
        elif current_triple:
            if current_count >= SUPPORT:
                print '%s\t%s\t%s' % (current_triple, current_count
, current_count / total)
            current_count = count
            current_triple = triple

# Output last triple
if current_triple == triple and current_count >= SUPPORT:
    print '%s\t%s\t%s' % (current_triple, current_count, current_co
unt / total)
```

Overwriting reducer_3_7_pass3.py

```
In [33]: !chmod a+x *_3_7_pass3.py
```

Obtain L3

```
In [39]: # Apriori
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_7_pass3
# Use 2 mappers and 2 reducers
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=2 \
-mapper $PWD/mapper_3_7_pass3.py \
-combiner $PWD/combiner_3_7_pass3.py \
-reducer $PWD/reducer_3_7_pass3.py \
-input /user/konniam/week_03/ProductPurchaseData.txt \
-output /user/konniam/week_03/hw_3_7_pass3 \
-file L1_counts.txt \
-file L2_counts.txt
```

16/02/03 23:58:38 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

16/02/03 23:58:39 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.

Deleted /user/konniam/week_03/hw_3_7_pass3

16/02/03 23:58:41 WARN streaming.StreamJob: -file option is depreca
ted, please use generic option -files instead.

16/02/03 23:58:41 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

packageJobJar: [L1_counts.txt, L2_counts.txt, /var/folders/l8/h51_5
9852qscq403fs6q0xlh0000gn/T/hadoop-unjar2932304002032762554/] [] /v
ar/folders/l8/h51_59852qscq403fs6q0xlh0000gn/T/streamjob34340505426
11831197.jar tmpDir=null

16/02/03 23:58:43 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/02/03 23:58:43 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032

16/02/03 23:58:44 INFO mapred.FileInputFormat: Total input paths to
process : 1

16/02/03 23:58:45 INFO mapreduce.JobSubmitter: number of splits:2

16/02/03 23:58:45 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0034

16/02/03 23:58:45 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0034

16/02/03 23:58:45 INFO mapreduce.Job: The url to track the job: htt
p://Konniam-MacBook-Air.local:8088/proxy/application_1454487534358
_0034/

16/02/03 23:58:45 INFO mapreduce.Job: Running job: job_145448753435
8_0034

16/02/03 23:58:55 INFO mapreduce.Job: Job job_1454487534358_0034 ru
nning in uber mode : false


```

16/02/03 23:58:55 INFO mapreduce.Job: map 0% reduce 0%
16/02/03 23:59:11 INFO mapreduce.Job: map 18% reduce 0%
16/02/03 23:59:17 INFO mapreduce.Job: map 23% reduce 0%
16/02/03 23:59:20 INFO mapreduce.Job: map 25% reduce 0%
16/02/03 23:59:23 INFO mapreduce.Job: map 30% reduce 0%
16/02/03 23:59:26 INFO mapreduce.Job: map 33% reduce 0%
16/02/03 23:59:29 INFO mapreduce.Job: map 38% reduce 0%
16/02/03 23:59:35 INFO mapreduce.Job: map 43% reduce 0%
16/02/03 23:59:38 INFO mapreduce.Job: map 48% reduce 0%
16/02/03 23:59:41 INFO mapreduce.Job: map 51% reduce 0%
16/02/03 23:59:44 INFO mapreduce.Job: map 56% reduce 0%
16/02/03 23:59:47 INFO mapreduce.Job: map 59% reduce 0%
16/02/03 23:59:50 INFO mapreduce.Job: map 61% reduce 0%
16/02/03 23:59:51 INFO mapreduce.Job: map 78% reduce 0%
16/02/03 23:59:53 INFO mapreduce.Job: map 80% reduce 0%
16/02/04 00:00:02 INFO mapreduce.Job: map 83% reduce 0%
16/02/04 00:00:09 INFO mapreduce.Job: map 83% reduce 8%
16/02/04 00:00:10 INFO mapreduce.Job: map 100% reduce 8%
16/02/04 00:00:11 INFO mapreduce.Job: map 100% reduce 17%
16/02/04 00:00:12 INFO mapreduce.Job: map 100% reduce 100%
16/02/04 00:00:13 INFO mapreduce.Job: Job job_1454487534358_0034 co
mpleted successfully
16/02/04 00:00:13 INFO mapreduce.Job: Counters: 49

```

File System Counters

```

FILE: Number of bytes read=437184
FILE: Number of bytes written=1352182
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3462103
HDFS: Number of bytes written=11138
HDFS: Number of read operations=12
HDFS: Number of large read operations=0
HDFS: Number of write operations=4

```

Job Counters

```

Launched map tasks=2
Launched reduce tasks=2
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)
=126742
Total time spent by all reduces in occupied slots (
ms)=33242
Total time spent by all map tasks (ms)=126742
Total time spent by all reduce tasks (ms)=33242
Total vcore-seconds taken by all map tasks=126742
Total vcore-seconds taken by all reduce tasks=33242
Total megabyte-seconds taken by all map tasks=12978
3808
Total megabyte-seconds taken by all reduce tasks=34
039808

```

Map-Reduce Framework

Map input records=31101
Map output records=279049
Map output bytes=6599573
Map output materialized bytes=437196
Input split bytes=236
Combine input records=279049
Combine output records=13871
Reduce input groups=7070
Reduce shuffle bytes=437196
Reduce input records=13871
Reduce output records=233
Spilled Records=27742
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=657
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=570949632

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=3461867

File Output Format Counters

Bytes Written=11138

16/02/04 00:00:13 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_7_pass3

```
In [40]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass3/part* | sort -k2,2nr | head
```

```
16/02/04 00:00:20 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable
```

DAI75645.FRO40251.SNA80324	550	0.0176843188322
DAI62779.FRO40251.SNA80324	476	0.0153049741166
FRO40251.GRO85051.SNA80324	471	0.0151442075817
DAI62779.ELE92920.SNA18336	432	0.01389022861
DAI62779.DAI75645.SNA80324	421	0.0135365422334
DAI62779.ELE17451.SNA80324	417	0.0134079290055
DAI62779.DAI75645.FRO40251	412	0.0132471624707
DAI62779.ELE17451.FRO40251	406	0.0130542426289
DAI75645.FRO40251.GRO85051	395	0.0127005562522
DAI62779.FRO40251.GRO85051	381	0.0122504099547

4th MR job to obtain confidence

```
In [45]: %%writefile mapper_3_7_pass4.py
#!/usr/bin/env python
# Mapper for 3.7
# Author: Konniam Chan
# Apriori Algorithm
from __future__ import division
from itertools import combinations
import sys

# Load L2 dictionary
L2_counts = {}
with open("L2_counts.txt", "r") as f:
    for line in f:
        pair, L2_support_count = line.strip().split()
        L2_counts[pair] = int(L2_support_count)

for line in sys.stdin:
    triple, L3_support_count, L3_support = line.strip().split()
    items = triple.split('.')
    # Obtain confidence
    for pair in combinations(items, 2):
        item_1, item_2 = pair[0], pair[1]
        item_3 = [x for x in items if x != item_1 and x != item_2][
0]

        L2_support_count = L2_counts['.'.join([item_1, item_2])]
        confidence = int(L3_support_count) / L2_support_count

        print ('%s.%s\t-> %s, %s,\t%s,\t%s' %
                (item_1, item_2, item_3, L3_support_count, L3_support
t, confidence))
```

Overwriting mapper_3_7_pass4.py

```
In [46]: %%writefile reducer_3_7_pass4.py
#!/usr/bin/env python
# Mapper for 3.7
# Author: Konniam Chan
# Apriori Algorithm
import sys

counter = 0
N = 10
for line in sys.stdin:
    # Output 10 lines only
    if counter >= N:
        continue
    print line.strip()
    counter += 1
```

Overwriting reducer_3_7_pass4.py

```
In [43]: !chmod a+x *_3_7_pass4.py
```

```
In [49]: # Apriori
!hdfs dfs -rm -r /user/konniam/week_03/hw_3_7_pass4
# Use 2 mappers and 1 reducer
!hadoop jar $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-*.
jar \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapr
educe.lib.partition.KeyFieldBasedComparator \
-D stream.num.map.output.key.fields=4 \
-D mapreduce.partition.keycomparator.options="-k4,4nr -k3,3nr" \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=1 \
-mapper $PWD/mapper_3_7_pass4.py \
-reducer $PWD/reducer_3_7_pass4.py \
-input /user/konniam/week_03/hw_3_7_pass3/part* \
-output /user/konniam/week_03/hw_3_7_pass4 \
-file L2_counts.txt
```

16/02/04 00:06:34 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

16/02/04 00:06:35 INFO fs.TrashPolicyDefault: Namenode trash config
uration: Deletion interval = 0 minutes, Emptier interval = 0 minute
s.

Deleted /user/konniam/week_03/hw_3_7_pass4

16/02/04 00:06:36 WARN streaming.StreamJob: -file option is depreca
ted, please use generic option -files instead.

16/02/04 00:06:36 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes whe
re applicable

```

packageJobJar: [L2_counts.txt, /var/folders/l8/h5l_59852qscq403fs6q
0x1h0000gn/T/hadoop-unjar3104935187276488070/] [] /var/folders/l8/h
5l_59852qscq403fs6q0x1h0000gn/T/streamjob102885312347332334.jar tmp
Dir=null
16/02/04 00:06:37 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/04 00:06:38 INFO client.RMProxy: Connecting to ResourceManage
r at /0.0.0.0:8032
16/02/04 00:06:39 INFO mapred.FileInputFormat: Total input paths to
process : 2
16/02/04 00:06:39 INFO mapreduce.JobSubmitter: number of splits:3
16/02/04 00:06:39 INFO mapreduce.JobSubmitter: Submitting tokens fo
r job: job_1454487534358_0036
16/02/04 00:06:40 INFO impl.YarnClientImpl: Submitted application a
pplication_1454487534358_0036
16/02/04 00:06:40 INFO mapreduce.Job: The url to track the job: htt
p://Konniams-MacBook-Air.local:8088/proxy/application_1454487534358
_0036/
16/02/04 00:06:40 INFO mapreduce.Job: Running job: job_145448753435
8_0036
16/02/04 00:06:49 INFO mapreduce.Job: Job job_1454487534358_0036 ru
nning in uber mode : false
16/02/04 00:06:49 INFO mapreduce.Job:  map 0% reduce 0%
16/02/04 00:07:03 INFO mapreduce.Job:  map 67% reduce 0%
16/02/04 00:07:04 INFO mapreduce.Job:  map 100% reduce 0%
16/02/04 00:07:12 INFO mapreduce.Job:  map 100% reduce 100%
16/02/04 00:07:12 INFO mapreduce.Job: Job job_1454487534358_0036 co
mpleted successfully
16/02/04 00:07:12 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=49920
        FILE: Number of bytes written=576737
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=12761
        HDFS: Number of bytes written=577
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=3
        Launched reduce tasks=1
        Data-local map tasks=3
        Total time spent by all maps in occupied slots (ms)
=38243
        Total time spent by all reduces in occupied slots (
ms)=4924
        Total time spent by all map tasks (ms)=38243
        Total time spent by all reduce tasks (ms)=4924

```

```

Total vcore-seconds taken by all map tasks=38243
Total vcore-seconds taken by all reduce tasks=4924
Total megabyte-seconds taken by all map tasks=39160
832
Total megabyte-seconds taken by all reduce tasks=50
42176

```

Map-Reduce Framework

```

Map input records=233
Map output records=699
Map output bytes=48516
Map output materialized bytes=49932
Input split bytes=354
Combine input records=0
Combine output records=0
Reduce input groups=699
Reduce shuffle bytes=49932
Reduce input records=699
Reduce output records=10
Spilled Records=1398
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=665
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=688914432

```

Shuffle Errors

```

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

```

File Input Format Counters

```
Bytes Read=12407
```

File Output Format Counters

```
Bytes Written=577
```

```

16/02/04 00:07:12 INFO streaming.StreamJob: Output directory: /user
/konniam/week_03/hw_3_7_pass4

```

Top 10 Rules by Confidence

(sorted by support if tied)

Results Formuat

(item1, item2) -> item3, supportCount, support, confidence

```
In [50]: !hdfs dfs -cat /user/konniam/week_03/hw_3_7_pass4/part*
```

```
16/02/04 00:07:17 WARN util.NativeCodeLoader: Unable to load native
-hadoop library for your platform... using builtin-java classes where applicable
```

```
GRO85051.SNA80324      -> FRO40251, 471,      0.0151442075817,
1.0
DAI75645.GRO85051      -> FRO40251, 395,      0.0127005562522,
1.0
ELE17451.GRO85051      -> FRO40251, 217,      0.00697726761197,
1.0
GRO73461.GRO85051      -> FRO40251, 147,      0.00472653612424,
1.0
ELE26917.GRO85051      -> FRO40251, 146,      0.00469438281727,
1.0
DAI23334.ELE92920      -> DAI62779, 143,      0.00459792289637,
1.0
ELE20847.GRO85051      -> FRO40251, 139,      0.0044693096685,
1.0
DAI55911.GRO85051      -> FRO40251, 133,      0.00427638982669,
1.0
ELE20847.FRO92469      -> FRO40251, 122,      0.00392270345005,
1.0
GRO21487.GRO85051      -> FRO40251, 120,      0.00385839683611,
1.0
```

HW3.8

Benchmark your results using the pyFIM implementation of the Apriori algorithm (Apriori - Association Rule Induction / Frequent Item Set Mining implemented by Christian Borgelt). You can download pyFIM from here:

<http://www.borgelt.net/pyfim.html> (<http://www.borgelt.net/pyfim.html>)

Comment on the results from both implementations (your Hadoop MapReduce of apriori versus pyFIM) in terms of results and execution times.

```
In [51]: from fim import apriori
```



```
In [57]: def apriori_3_8():
# Pre-process to create iterable
transactions = []
with open('ProductPurchaseData.txt', 'r') as f:
    for line in f:
        items = line.strip().split()
        transactions.append(items)

# Apriori
results = apriori(transactions, target='r', supp=-100, zmin=3,
zmax=3, report='asc')
# Sort by confidence, then by support
results = sorted(results, key=lambda x: (-x[4], -x[3]))
for row in results[:10]:
    print row[1], row[0], row[2], row[3], row[4]

apriori_3_8()

('GRO85051', 'SNA80324') FRO40251 471 0.0151442075817 1.0
('GRO85051', 'DAI75645') FRO40251 395 0.0127005562522 1.0
('GRO85051', 'ELE17451') FRO40251 217 0.00697726761197 1.0
('GRO85051', 'GRO73461') FRO40251 147 0.00472653612424 1.0
('GRO85051', 'ELE26917') FRO40251 146 0.00469438281727 1.0
('DAI23334', 'ELE92920') DAI62779 143 0.00459792289637 1.0
('ELE20847', 'GRO85051') FRO40251 139 0.0044693096685 1.0
('DAI55911', 'GRO85051') FRO40251 133 0.00427638982669 1.0
('ELE20847', 'FRO92469') FRO40251 122 0.00392270345005 1.0
('GRO85051', 'GRO21487') FRO40251 120 0.00385839683611 1.0
```

The results of Mapreduce and pyFIM implementations are identical. In terms of execution times, pyFIM took less than 1 second, while MR took about 3 mins in total. Because the data set is small in this case and could be held in memory effectively, the in-memory approach of pyFIM is much faster. However, if the data set were 1TB, Mapreduce could still handle it with multiple machines, whereas pyFIM will fail to work.